


File System Implementation

**Silberschatz, Galvin
Jul-Nov 2018
Moumita Patra**

- **File system**- provides the mechanism for on-line storage and access to file contents, including data and programs
- File system resides in secondary storage
 - Provide user interface to storage, mapping logical to physical
 - File systems provide efficient and convenient access to disk
- Problems to handle:
 - How the file system will look to the user
 - Map the logical file system onto the physical secondary storage devices

- **File Control Block**- storage structure consisting of information about a file
- **Device driver** controls the physical device
- File system organized into layers

Layered File System



Layered File System

- **I/O Control**- consists of device drivers and interrupt handlers to transfer information between the main memory and the disk system
- **Basic file system**- issues generic commands to the appropriate device driver to read and write physical blocks on the disk
- **File organization module**- understands files, logical address and physical blocks
- Translates logical block # to physical block #
- Manages free space, disk allocation

- **Logical file system**- manages metadata information
- Translates file name into file number, file handle , location by maintaining FCBs
- Directory management
- Protection

Layered sturture for file system helps in-

- Duplication of code is minimized
- Layering may introduce more OS overhead
- Challenges- how many layers to use? what each layer should do?

File-System Implementation

- We have system calls at the API level, but how do we implement their functions?
 - On-disk and in-memory structures
- **Boot control block** contains info needed by system to boot OS from that volume
 - Needed if volume contains OS, usually first block of volume
- **Volume control block (superblock, master file table)** contains volume details
 - Total # of blocks, # of free blocks, block size, free block pointers or array
- Directory structure organizes the files
 - Names and inode numbers, master file table

File-System Implementation (Cont.)

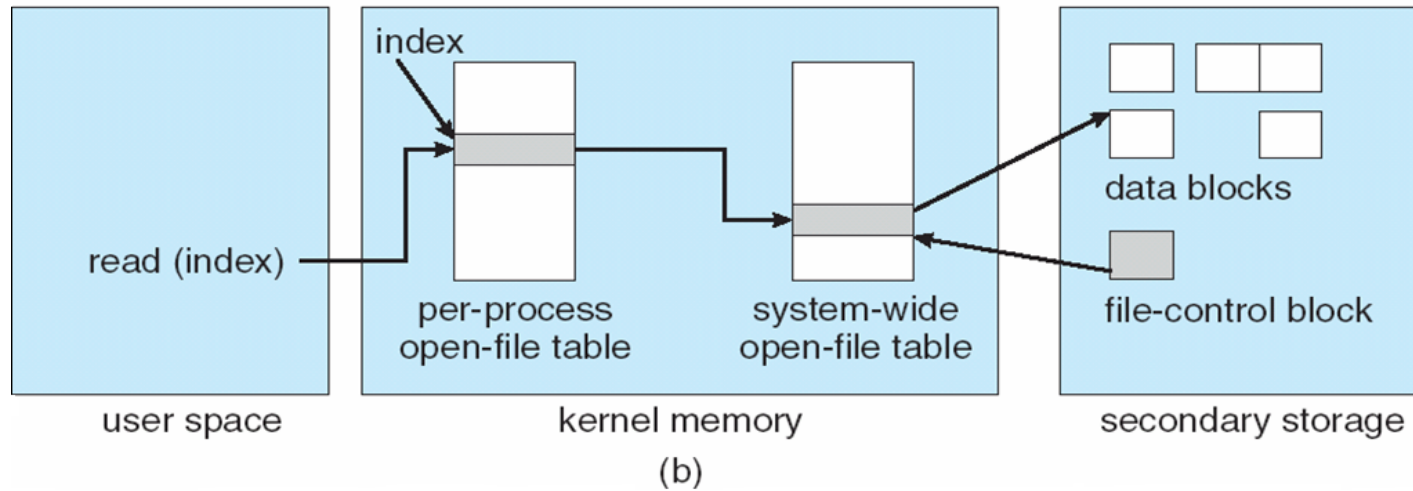

- Per-file **File Control Block (FCB)** contains many details about the file
 - inode number, permissions, size, dates
 - NTFS stores info in master file table using relational DB structures

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

In-Memory File Information

- In-memory information- used for file system management and performance improvement via caching
- In-memory structures:
- **Mount table**- contains information about each mounted volume
- **In-memory directory structure cache**- holds directory info of recently accessed directories
- **System-wide open file table**- contains a copy of the FCB of each open file as well as other info
- **Per-process open-file table**- contains a pointer to the appropriate entry in the system-wide open-file table
- **Buffers**- hold file-system blocks when they are being read/written from/to disks.

In-Memory File System Structures



Partitions and Mounting


- Partition can be a volume containing a file system (“cooked”) or **raw** – just a sequence of blocks with no file system
- Boot block can point to boot volume or boot loader set of blocks that contain enough code to know how to load the kernel from the file system
 - Or a boot management program for multi-os booting
- **Root partition** contains the OS, other partitions can hold other Oses, other file systems, or be raw
 - Mounted at boot time
 - Other partitions can mount automatically or manually
- At mount time, file system consistency checked
 - Is all metadata correct?
 - If not, fix it, try again
 - If yes, add to mount table, allow access

Allocation Methods - Contiguous

- An allocation method refers to how disk blocks are allocated for files:
- **Contiguous allocation** – each file occupies set of contiguous blocks
 - Best performance in most cases
 - Simple – only starting location (block #) and length (number of blocks) are required
 - Problems include finding space for file, knowing file size, external fragmentation, need for **compaction off-line (downtime)** or **on-line**

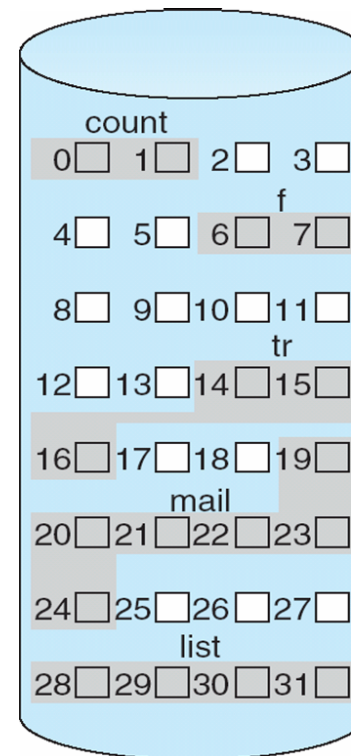
Contiguous Allocation

- Mapping from logical to physical



Block to be accessed = $Q + \text{starting address}$

Displacement into block = R



directory


file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

Allocation Methods - Linked


- **Linked allocation** – each file a linked list of blocks
 - File ends at null pointer
 - No external fragmentation
 - Each block contains pointer to next block
 - No compaction, external fragmentation
 - Free space management system called when new block needed
 - Improve efficiency by clustering blocks into groups but increases internal fragmentation
 - Reliability can be a problem
 - Locating a block can take many I/Os and disk seeks

Linked Allocation

- Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk




■ Mapping



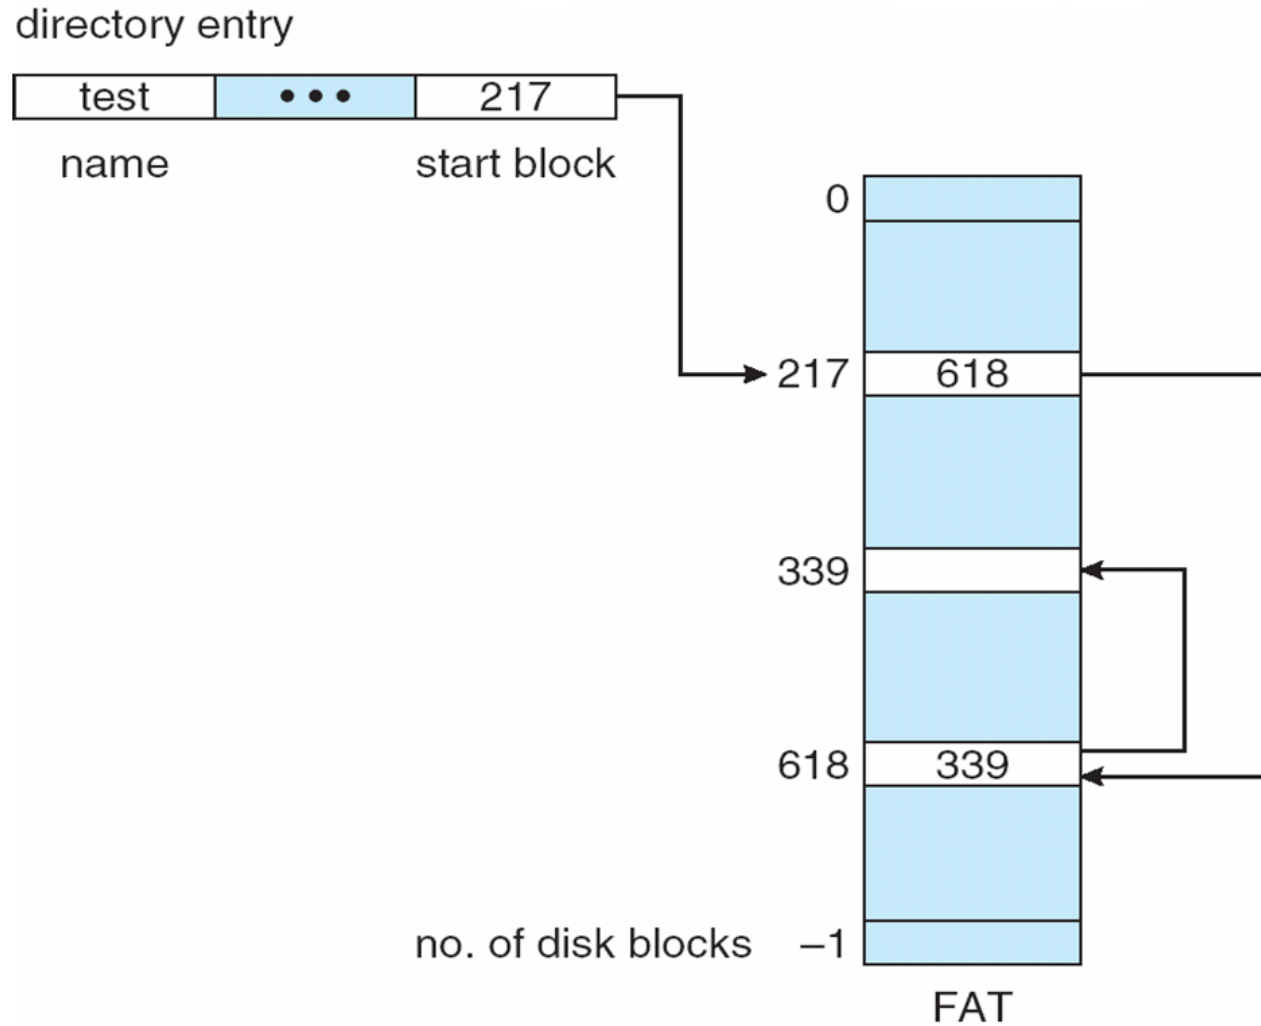
Block to be accessed is the Qth block in the linked chain of blocks representing the file.

Displacement into block = $R + 1$

Linked Allocation




File-Allocation Table



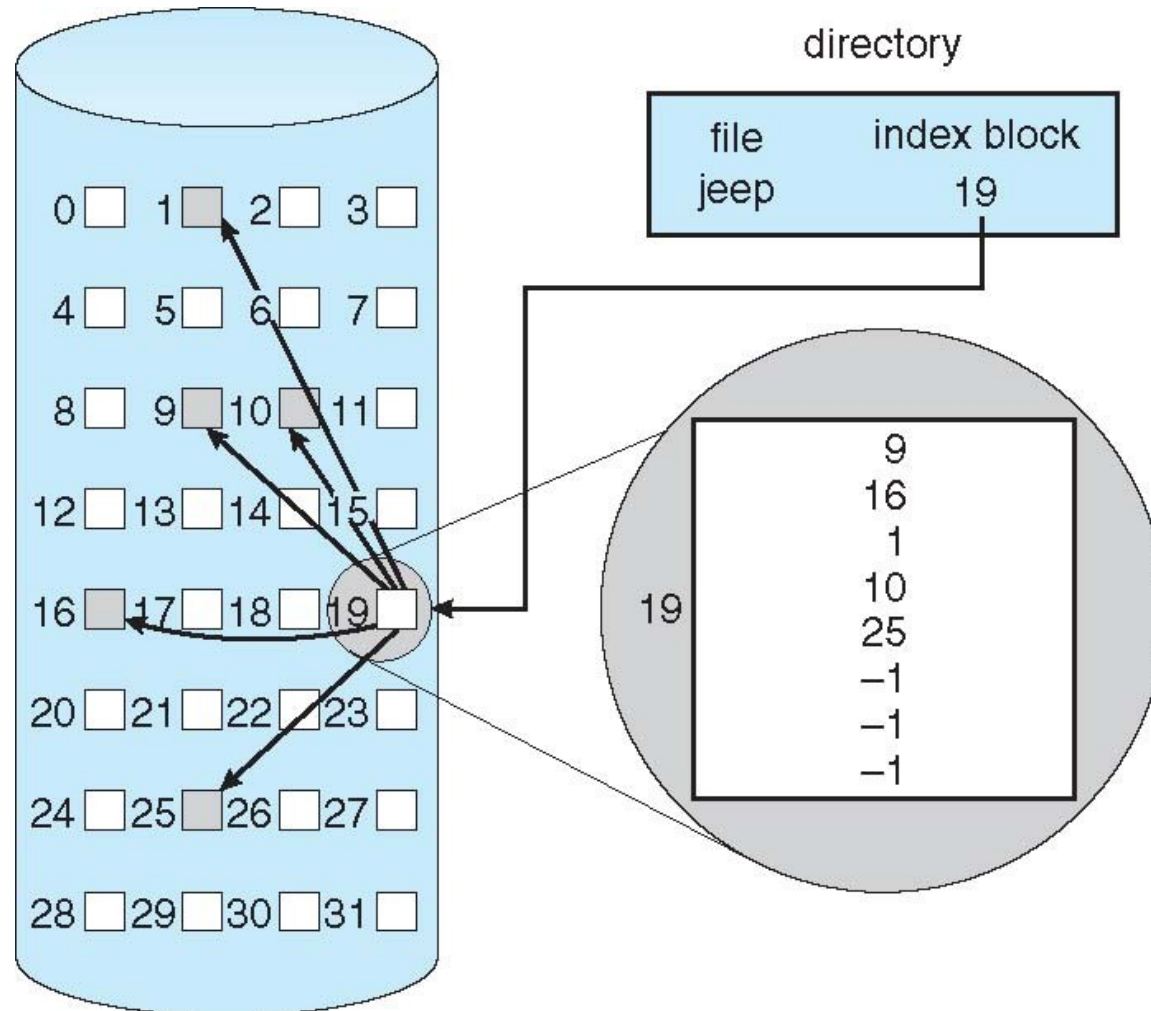
Allocation Methods - Indexed

- **Indexed allocation**
 - Each file has its own **index block**(s) of pointers to its data blocks

- Logical view




Example of Indexed Allocation



Indexed Allocation (Cont.)

- Need index table
- Random access
- Dynamic access without external fragmentation, but have overhead of index block
- Mapping from logical to physical in a file of maximum size of 256K bytes and block size of 512 bytes. We need only 1 block for index table



Q = displacement into index table

R = displacement into block