

Process Concept

Moumita Patra
July-Nov 2019
Ref: Galvin- Gagne


What is a process?

- **Process**- it is an instance of a program in execution.
- Multiple instances of the same program are different processes
- Process execution must progress in sequential fashion
- Process has resources allocated to it by the OS during execution
 - CPU time
 - Memory space for code, data, stack
 - Open files
 - Signals, etc.

- Each process is identified by a unique, positive integer id called **process id**.
- **Program-** is a passive entity stored on disk (**executable file**)
- Process is an **active** entity
- Program becomes process when executable file is loaded into memory
- One program can be several processes

Process- more than program code !

- Has multiple parts
 - **Text section**- the program code
 - **Stack**- contains temporary data (function parameters, return addresses, local variables, etc)
 - **Data section**- contains global variables
 - **Heap**- dynamically allocated memory during run time

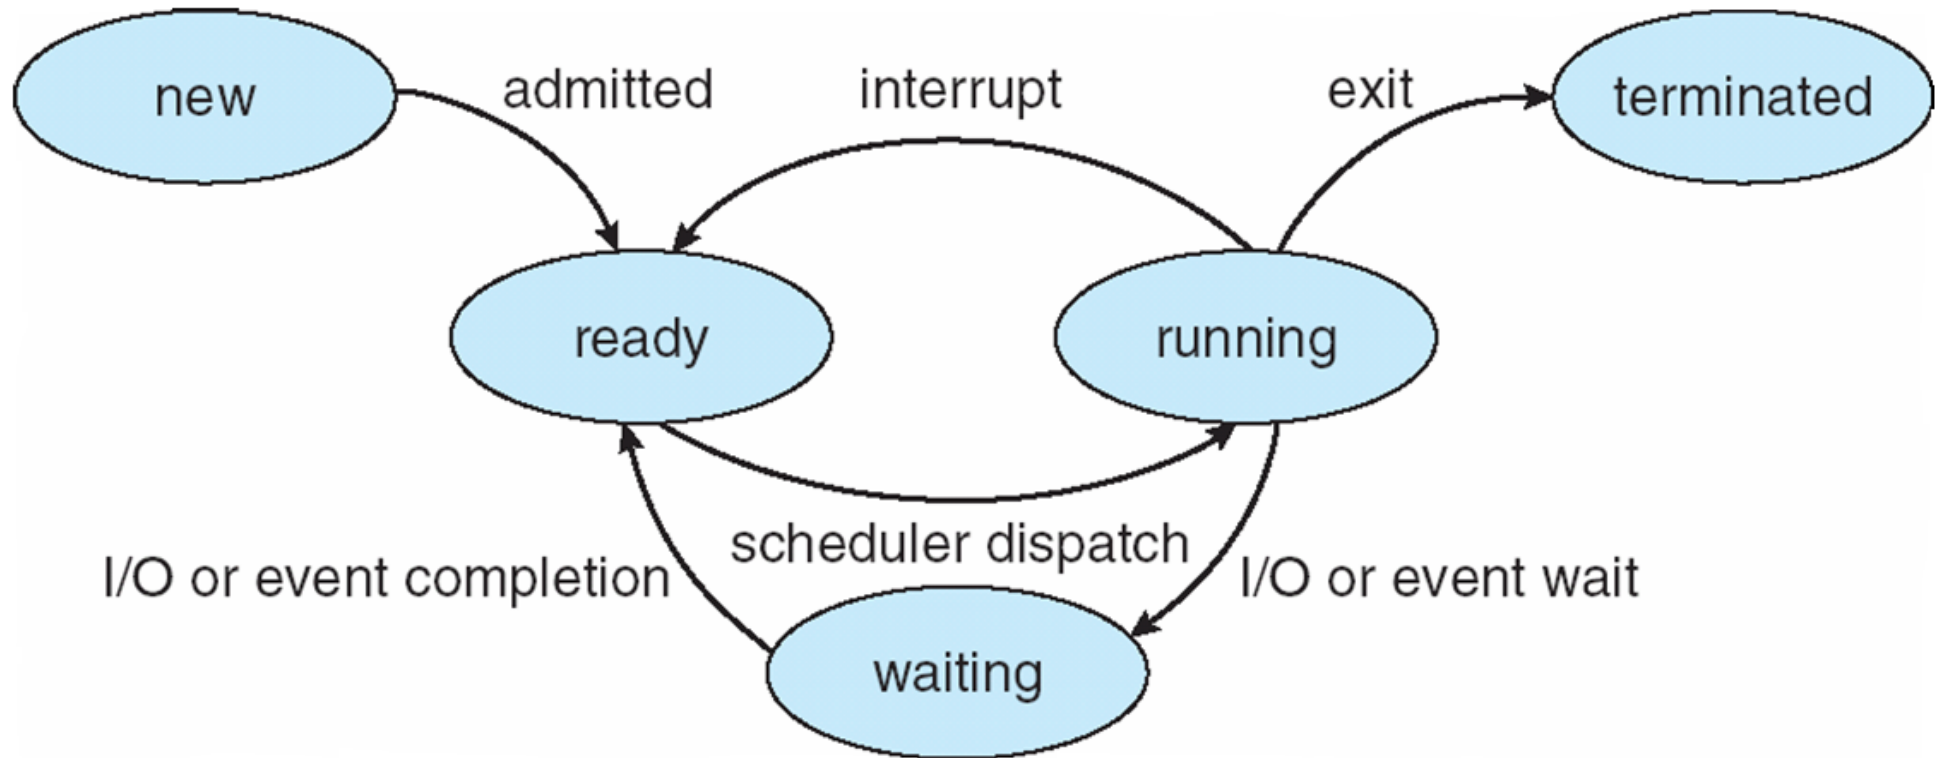


Process States

As a process executes, it changes **state**.

- **New-** The process is being created
- **Running-** Instructions are being executed
- **Waiting-** The process is waiting for some event to occur
- **Ready-** The process is waiting to be assigned to a processor
- **Terminated-** The process has finished execution

Process States

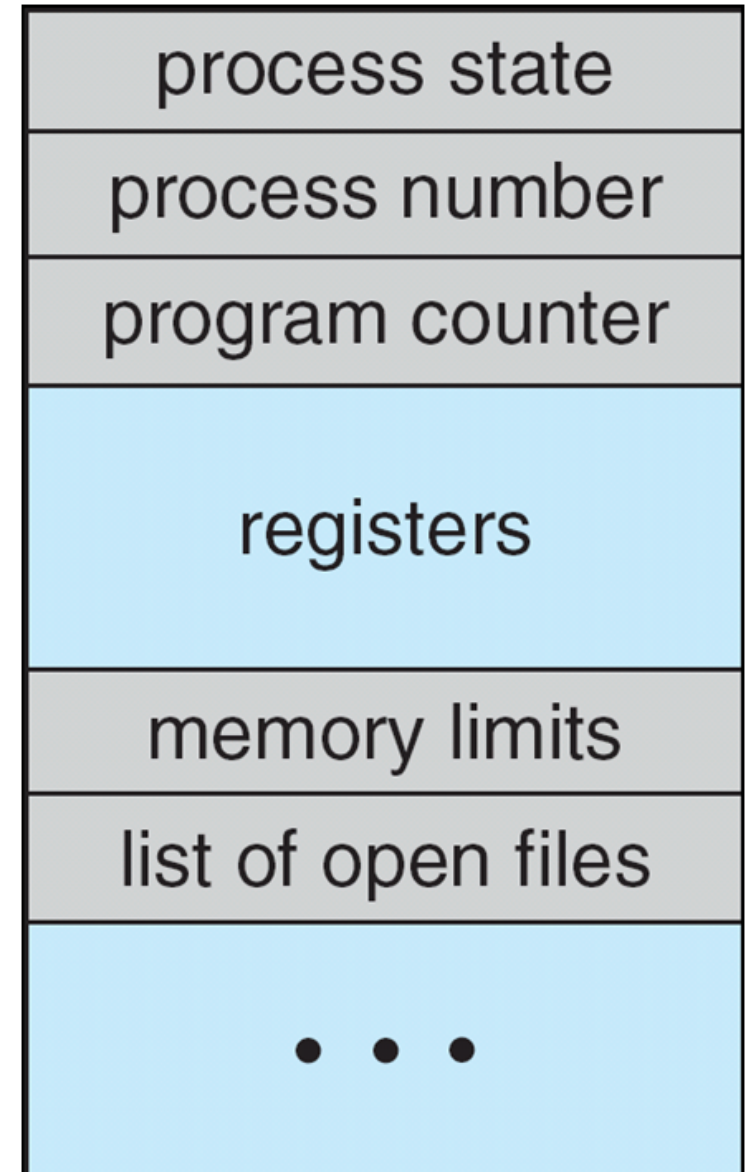


Process Control Block (PCB)

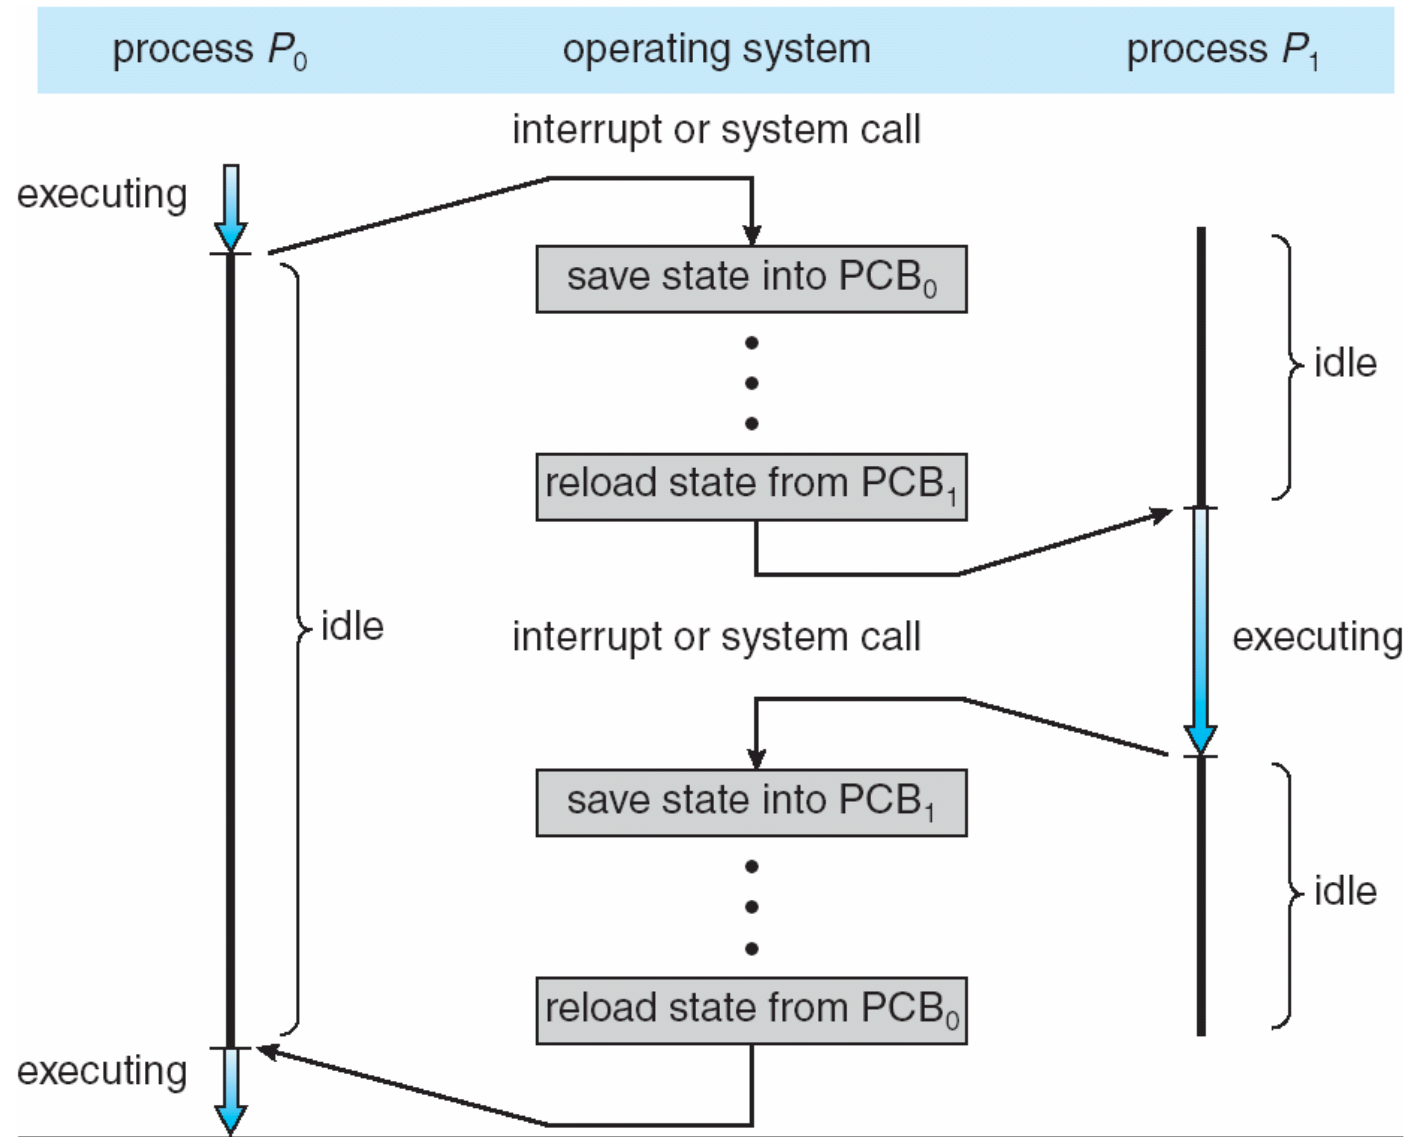
- Primary data structure maintained by the OS that contains information about a process
- One PCB per process
- OS maintains a list of PCBs for all processes
- **Interrupt**- A signal to the processor emitted by the hardware or software indicating an event that needs immediate attention.
- Processor responds to an interrupt by suspending its current activities, saving its state, and executing a function called an **interrupt handler** (or an **Interrupt Service Routine, ISR**).

Typical contents of PCB

- **Process state**- running, waiting, etc
- **Program counter**- location of instruction to execute next
- **CPU registers**- contents of all process-centric registers
- **CPU scheduling information**
- **Memory management information**
- **Accounting information**- CPU used, time elapsed, etc
- **I/O status information**- I/O devices allocated to process, list of open files



CPU Switch from Process to Process



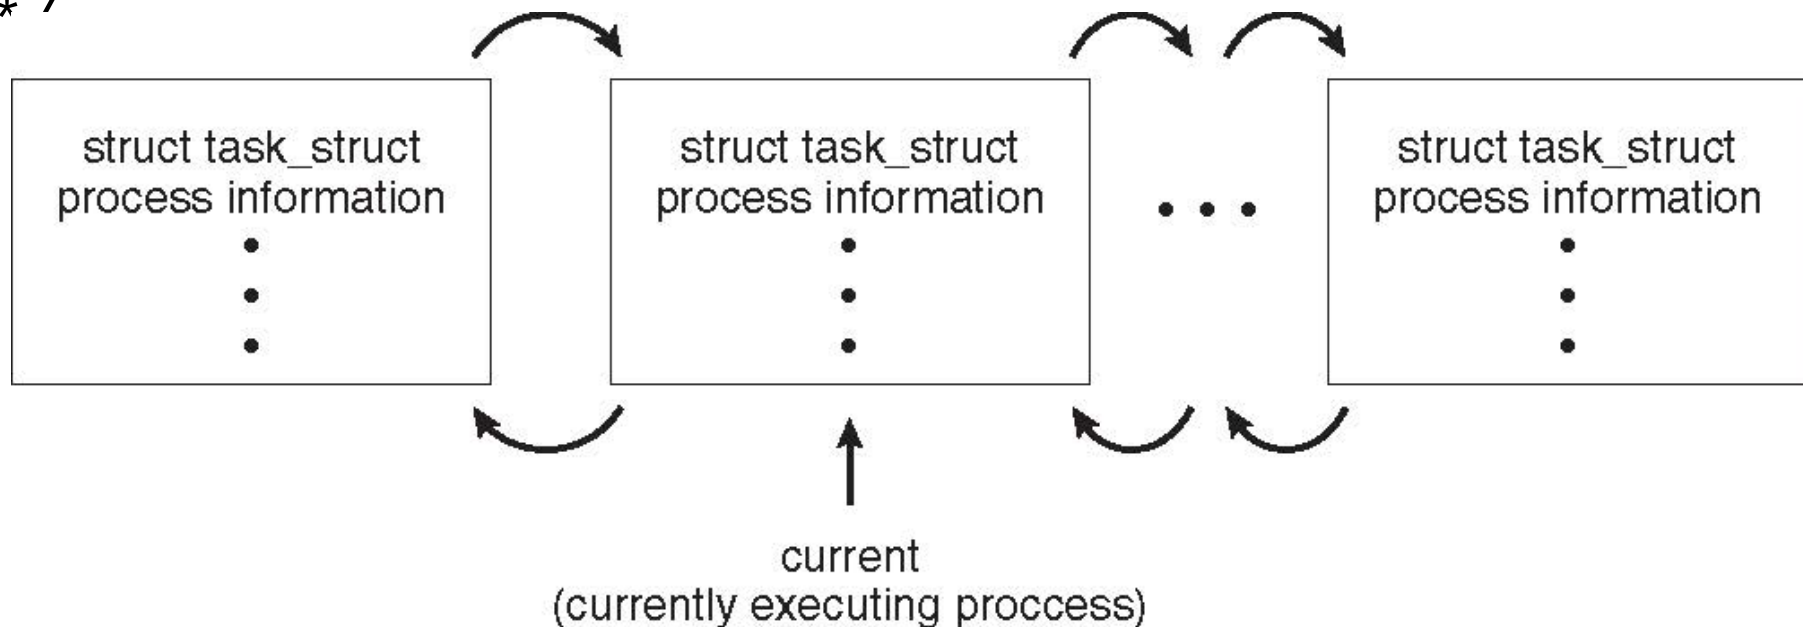
Threads

- A basic unit of CPU utilization.
- It is the smallest sequence of programmed instructions that can be executed independently.
- Traditionally, processes have a single thread of control.
- Multiple threads to perform more than one task at a time-- multicore systems.
- Therefore, need to have storage for thread details

Process Representation in Linux

Represented by the C structure `task_struct`

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children*/
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process
* '
```



Main Operations on a Process

- **Process creation**

- Data structures like PCB is set up and initialized
- Initial resources allocated and initialized if needed
- Process added to ready queue

- **Process scheduling**

- CPU is allotted to the process, process runs

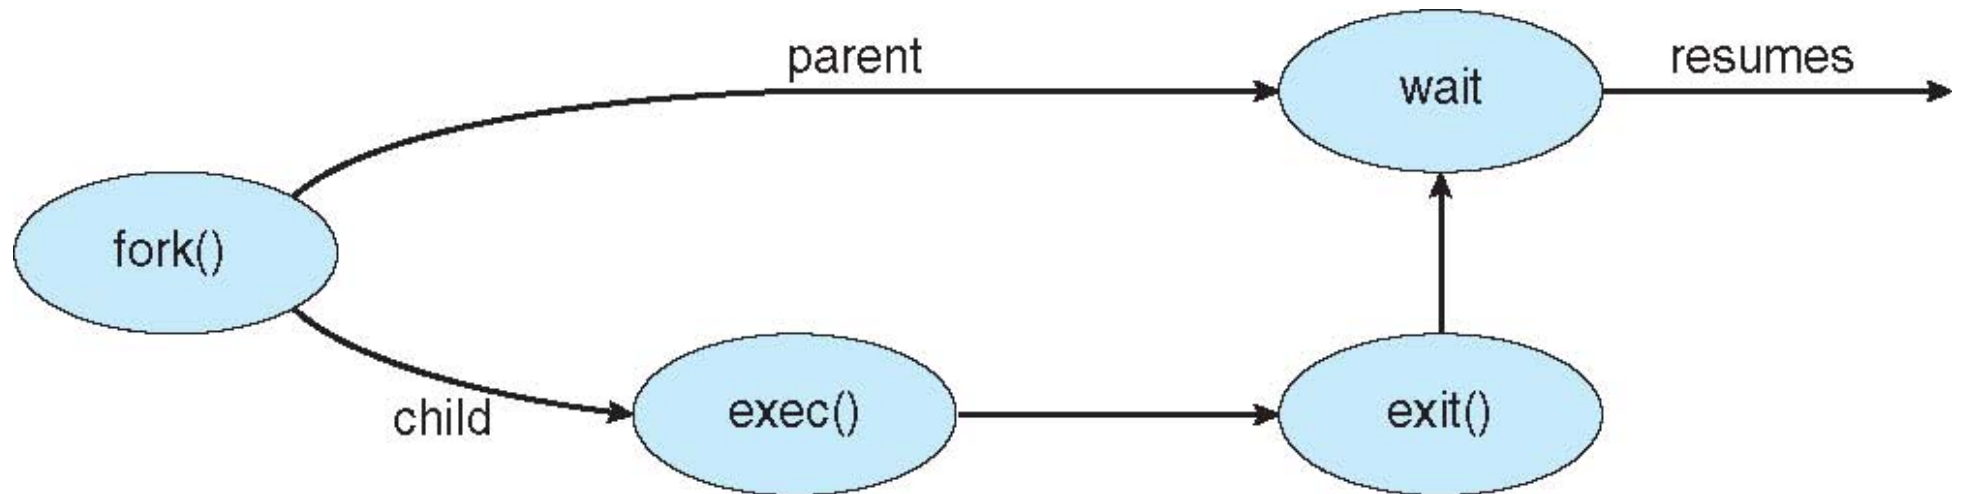
- **Process termination**

- Process is removed
- Resources are reclaimed

Process Creation

- A process (***parent***) can create another process (***child***), which in turn can create other processes, forming a ***tree*** of processes
 - By making a system call, ***fork()***
 - **Parent process**- process that invokes the call
 - **Child process**- the new process created
- Processes are identified and managed via ***process identifier (pid)***
- ***exec()*** system call used after a ***fork()*** to replace the process' memory space with a new program

Process Creation



C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

- **Resource sharing possibilities**

- Parent and children share all resources
- Children share subset of parent's resources
- Parent and child share no resources
- Restricting a child process to a subset of the parent's resources prevents any process from overloading the system by creating too many child processes.

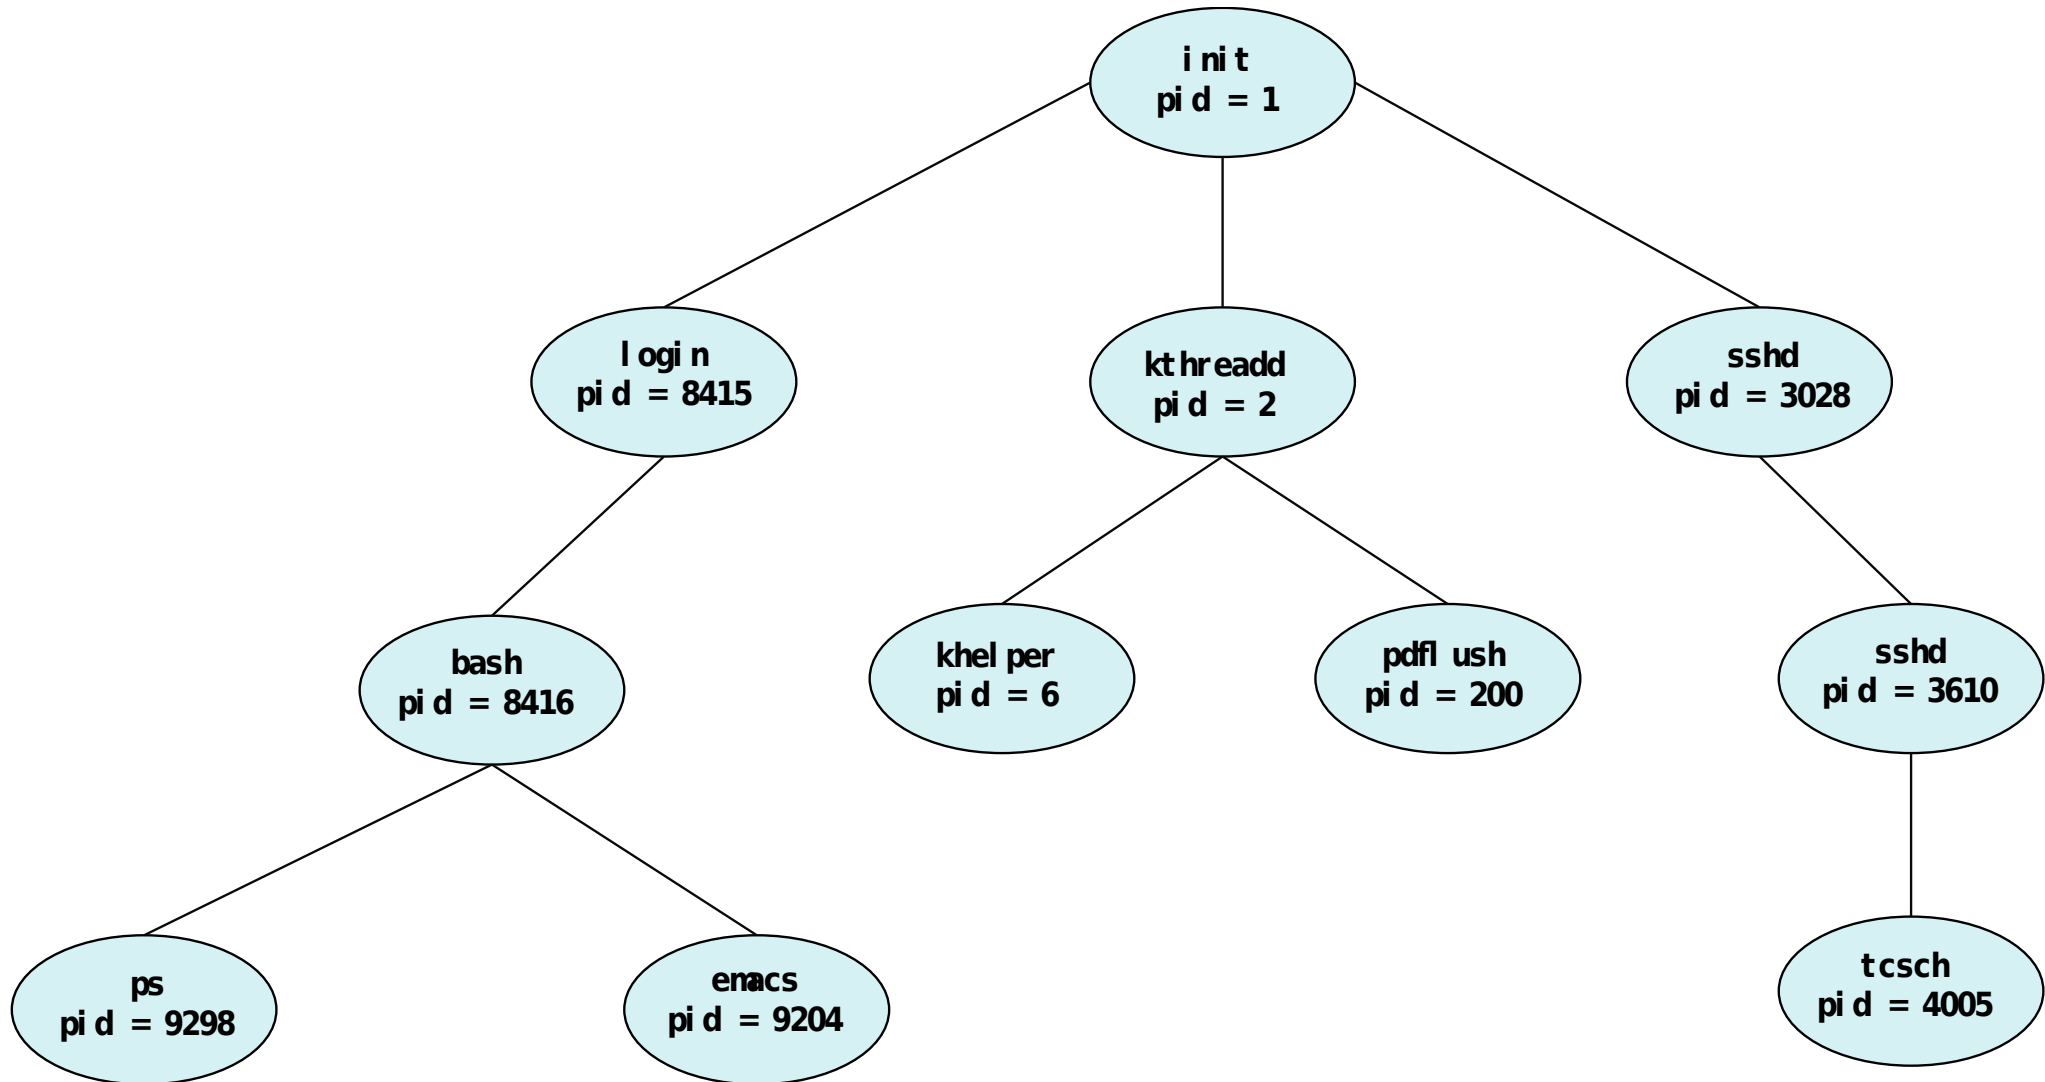
- **Execution possibilities**

- Parent and children execute concurrently
- Parent waits until children terminate

- **Memory address space possibilities**

- Address space of child duplicate of parent
- Child has a new program loaded into it

A Tree of Processes in Linux



Process Termination

- Process executes last statement and asks the OS to terminate it (e.g- ***exit()*** system call)
- Returns status data from child to parent (via ***wait()***)
- Process's resources are deallocated by OS
- Parent may terminate execution of a child using ***abort()*** due to-
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and OS does not allow a child to continue if its parent terminates.

Process Termination

- Some operating systems do not allow child to exist if its parent has terminated.
- If parent terminates then children must also terminate
- **Cascading termination**- all children, grandchildren, etc are terminated. Termination is initiated by OS.
- Parent process may wait for termination of a child process by using **wait()**. It returns status information and the pid of the terminated process.
- **Zombie process**- A process that has terminated but whose parent has not yet called **wait()**
- **Orphan process**- Parent terminated without **wait()**
- Linux and UNIX address this scenario by assigning **init()** process as the new parent to the orphan process

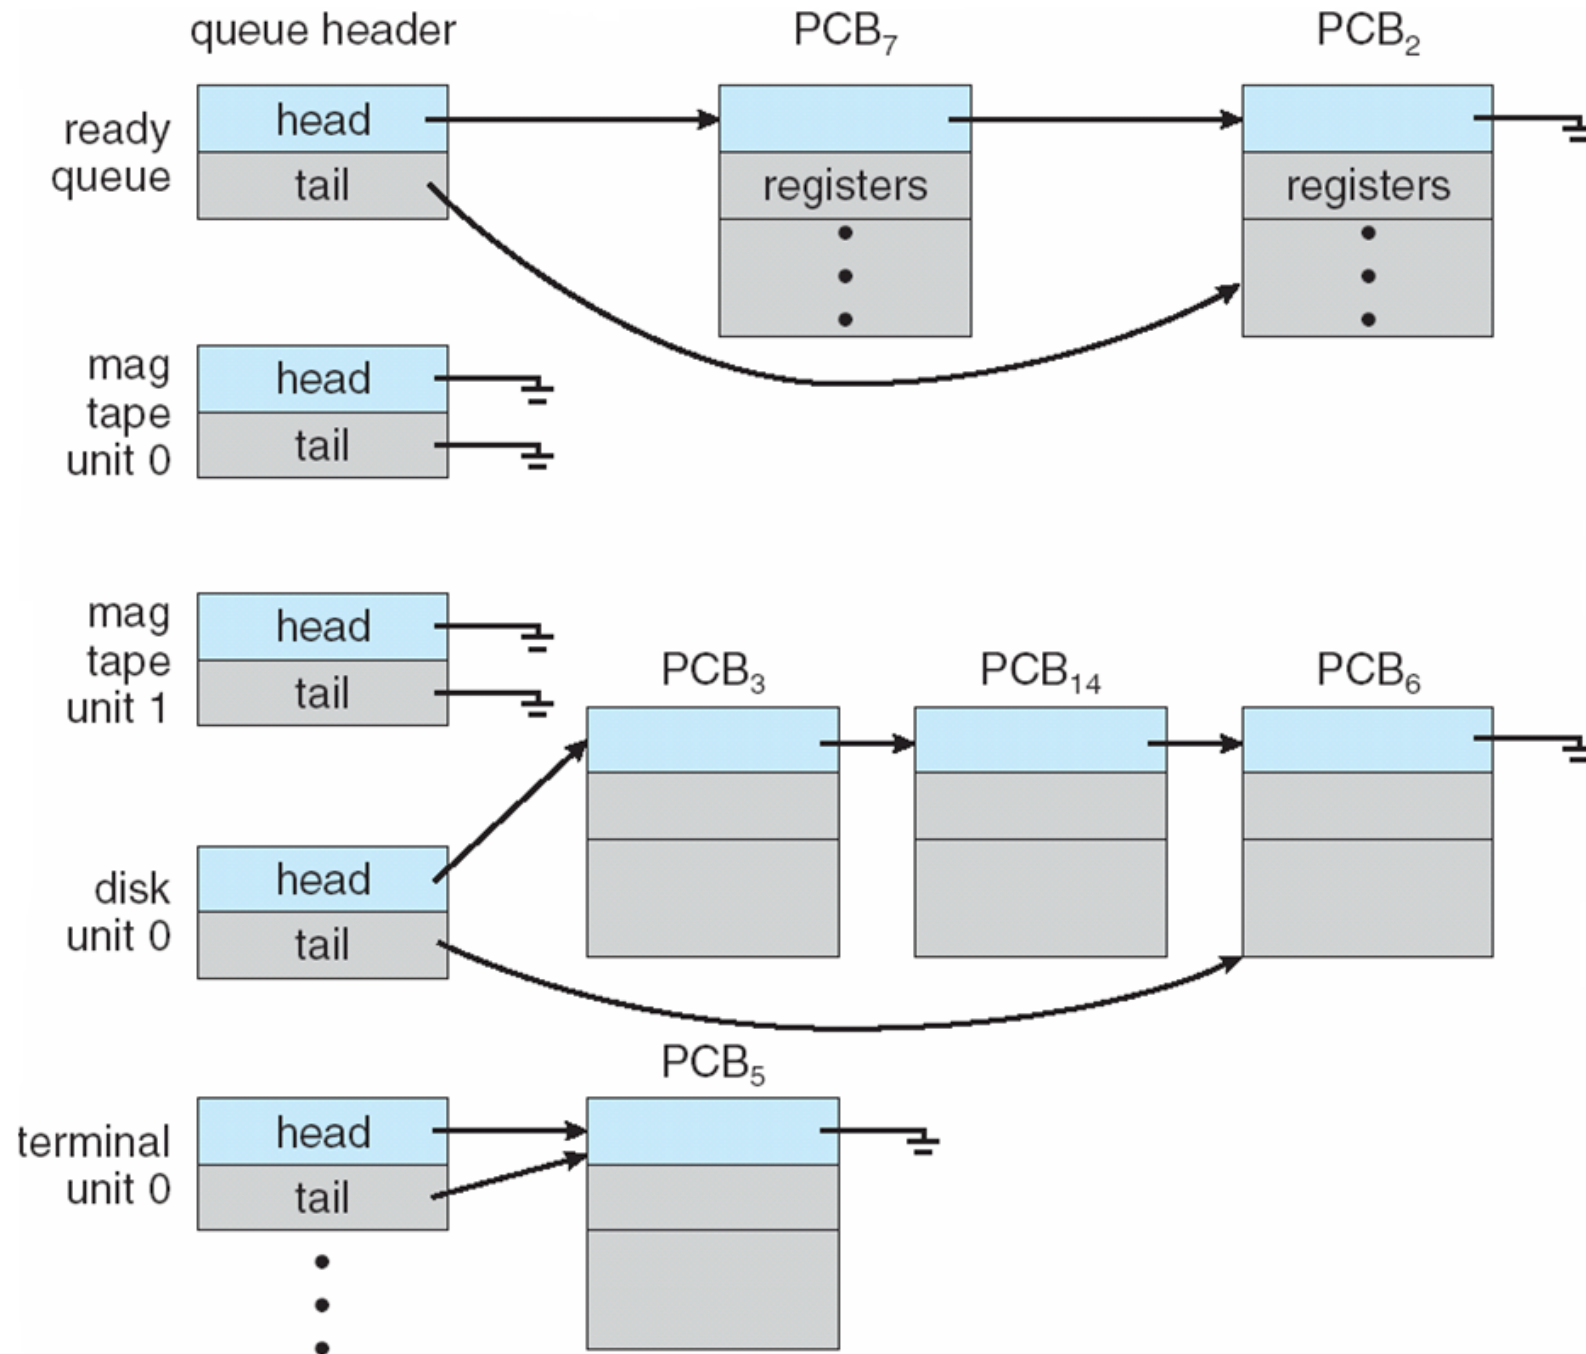
Process Scheduling

- Objective of multiprogramming is to have some process running at all times, to maximize CPU utilization.
- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler/ dispatcher** selects among available processes (from ready queue according to some algorithm) for next execution on CPU
- Selected process runs till
 - It needs to wait for some event to occur (e.g- disk read)
 - CPU time allotted expires
 - Arrival of a higher priority process

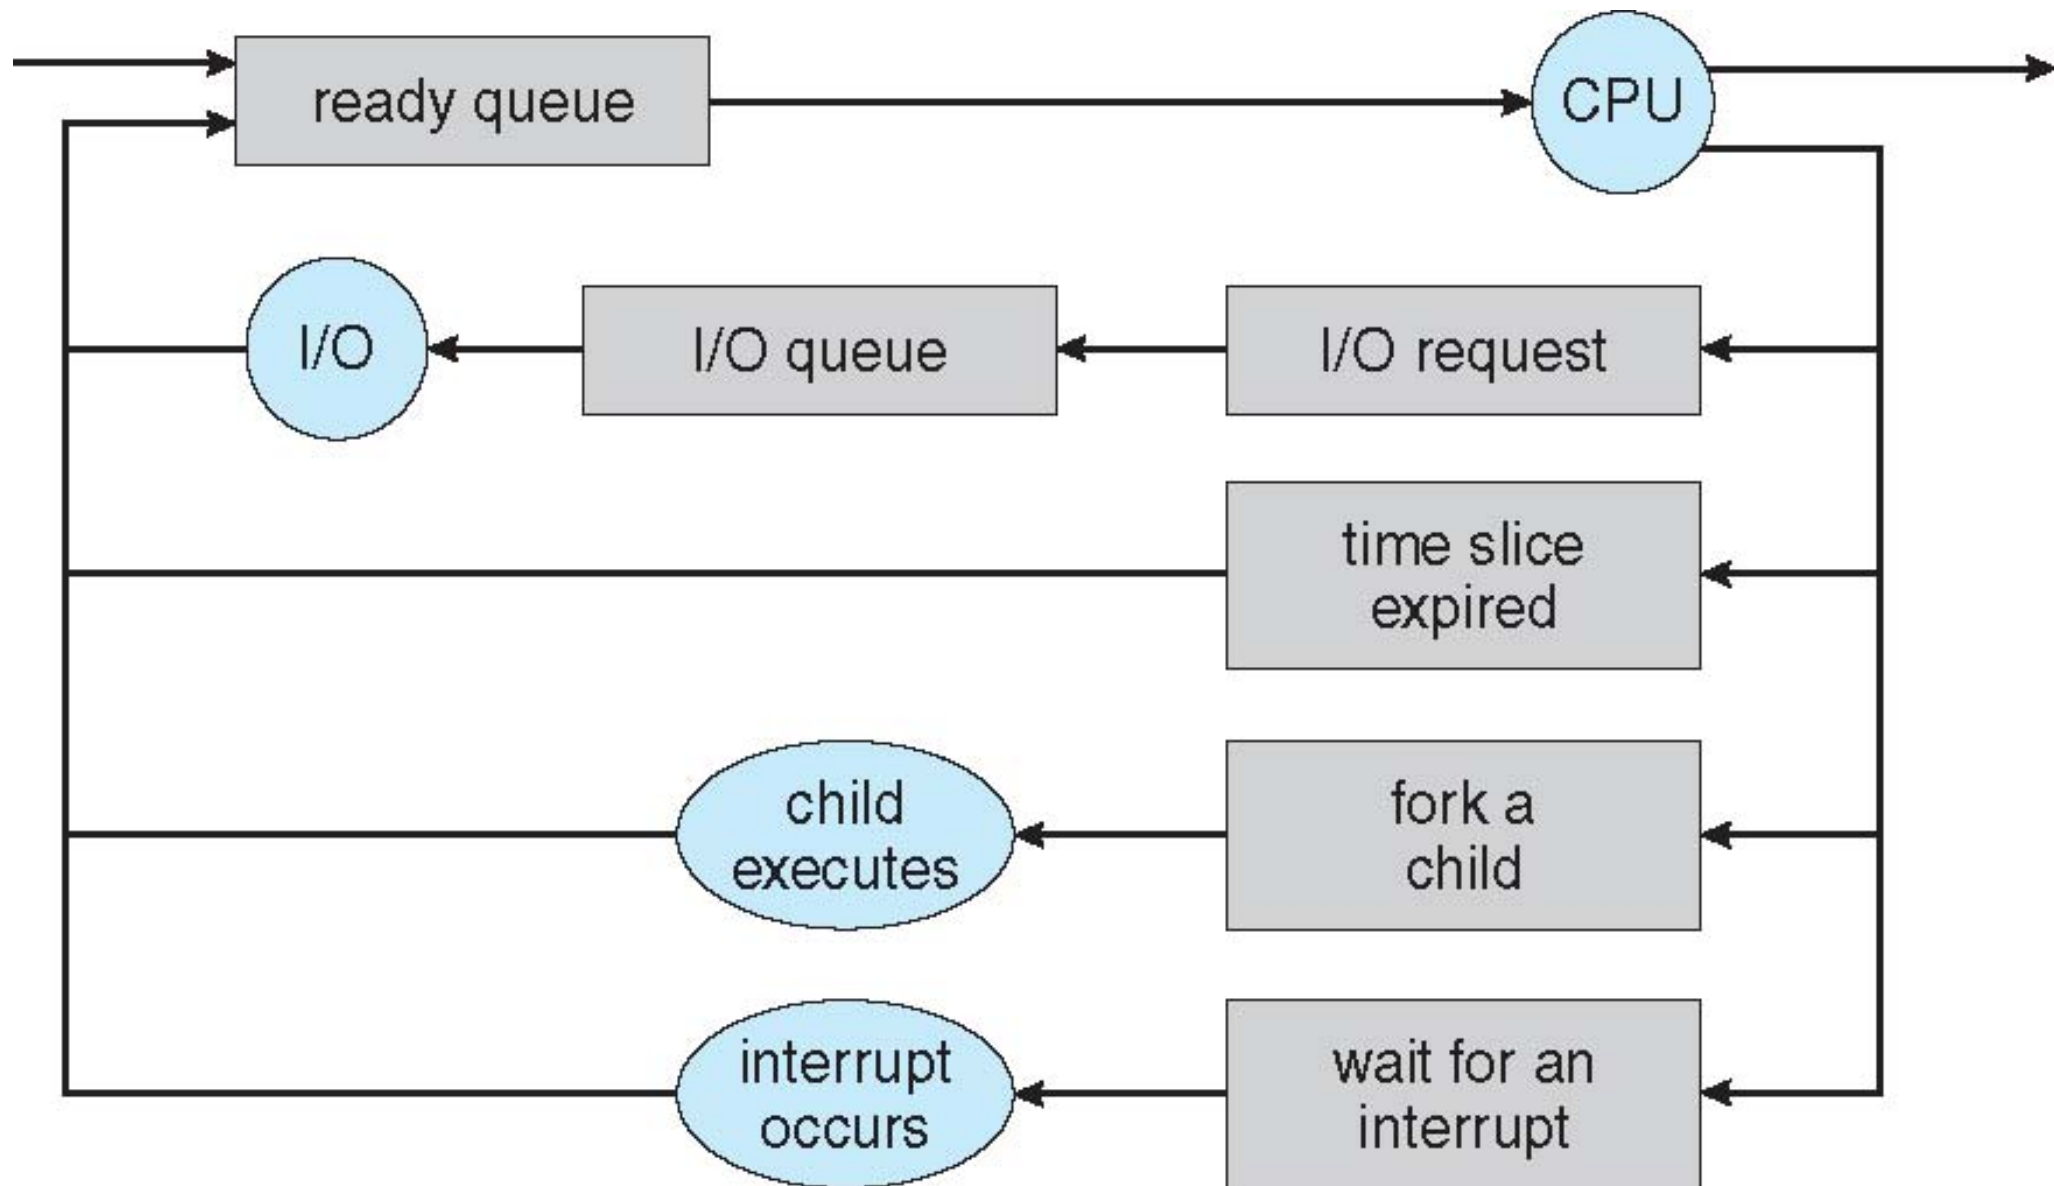
- When process is ready to run again, it goes back to the ready queue
- Scheduler is invoked again to select the next process from the ready queue

- Scheduling queues of processes-
 - Job queue- set of all processes in the system
 - Ready queue- set of all processes residing in main memory, ready and waiting to execute
 - Device queues- set of processes waiting for an I/O device
- Processes migrate among the various queues

Ready Queue and Various I/O Device Queues



Representation of Process Scheduling



Schedulers

- **Short term scheduler (or CPU scheduler)**- selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler must select a new process for the CPU frequently.
- **Long term scheduler (or job scheduler)**- selects which processes should be brought into the ready queue
 - Invoked less frequently.
 - Controls the **degree of multiprogramming** (the number of processes in memory)

Schedulers

Processes can be described as either:

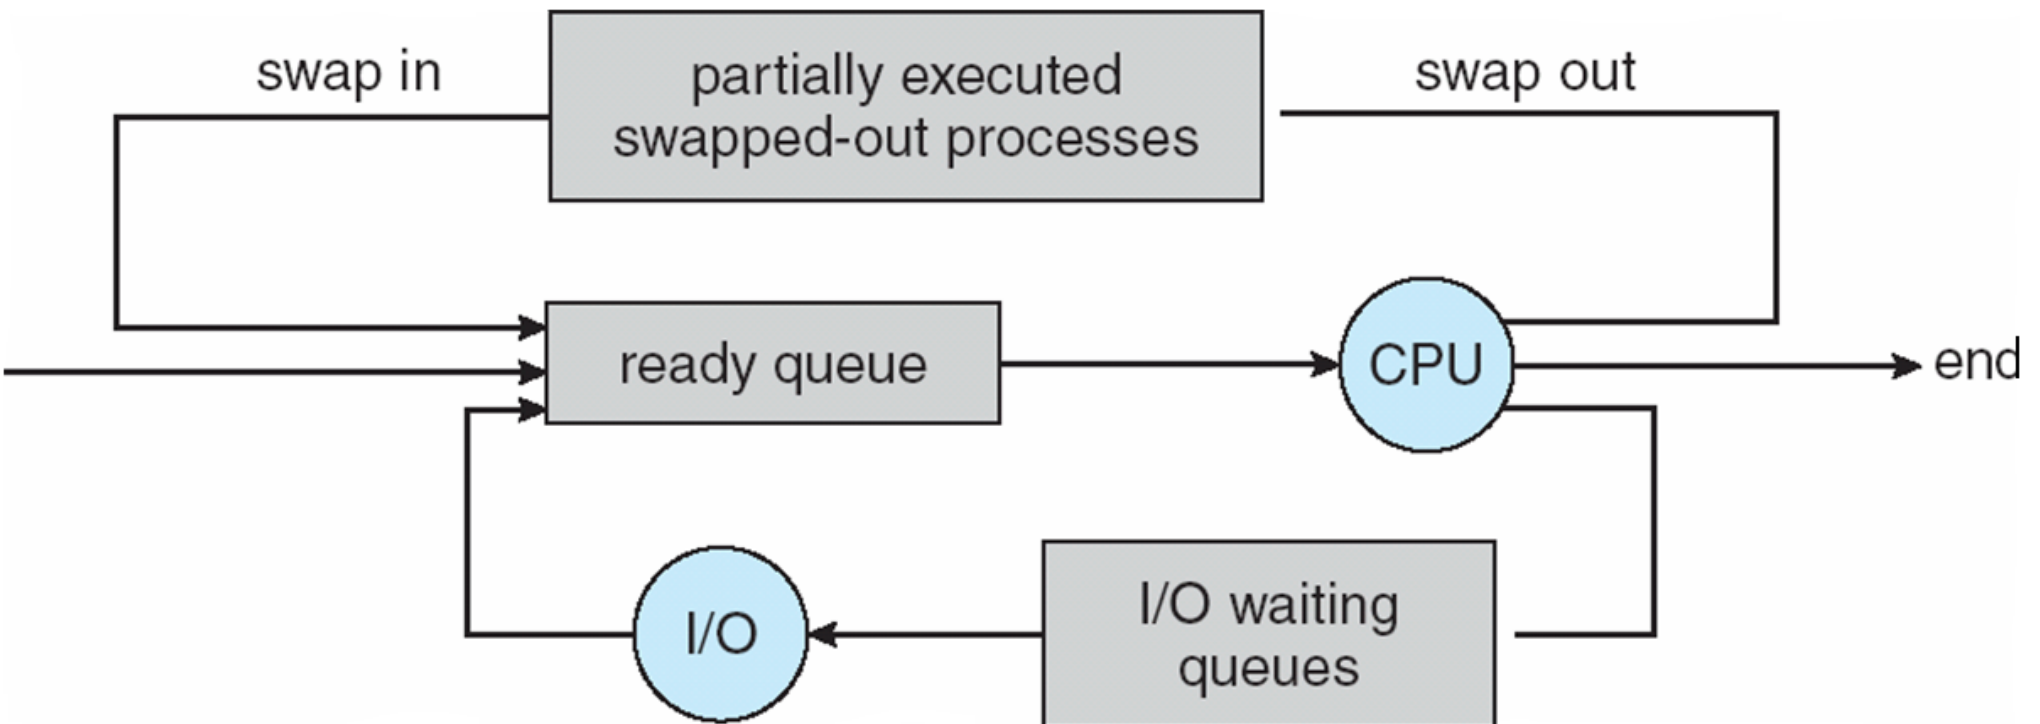
- **I/O bound process**- spends more time doing I/O than computations; many short CPU bursts
- **CPU-bound process**- spends more time doing computations; few very long CPU bursts
- Job of long term scheduler to make a careful selection.
- It is important that the long term scheduler selects a good mix of CPU and I/O bound processes.
- What if all processes are I/O bound?
- What if all processes are CPU bound?

- **Note:** In some systems, long time scheduler may be absent or minimal, e.g- UNIX and Windows. Every process is put into memory.
- Stability of such systems depends on physical limitations or on user's demand of work.
- **What if all processes do not fit in memory?**
 - Partially executed jobs in secondary memory (swapped out)
 - Copy the process image to some pre-designated area in the disk (swap out)
 - Bring in again later and add to ready queue later

Medium Term Scheduler

Some time-sharing systems may introduce additional, intermediate level of scheduling


- Helps in removing process from memory, store on disk, bring back in from disk to continue execution : swapping



Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the saved state for the new process via a **context switch**
- **Context**- Information that is required to be saved to be able to restart the process later from the same point
- **Context** of a process is represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
- More complex the OS and the PCB-> longer the context switch
- Switching time dependent on hardware support- memory speed, number of registers that must be copied, special instructions, etc.

Interprocess Communication(IPC)

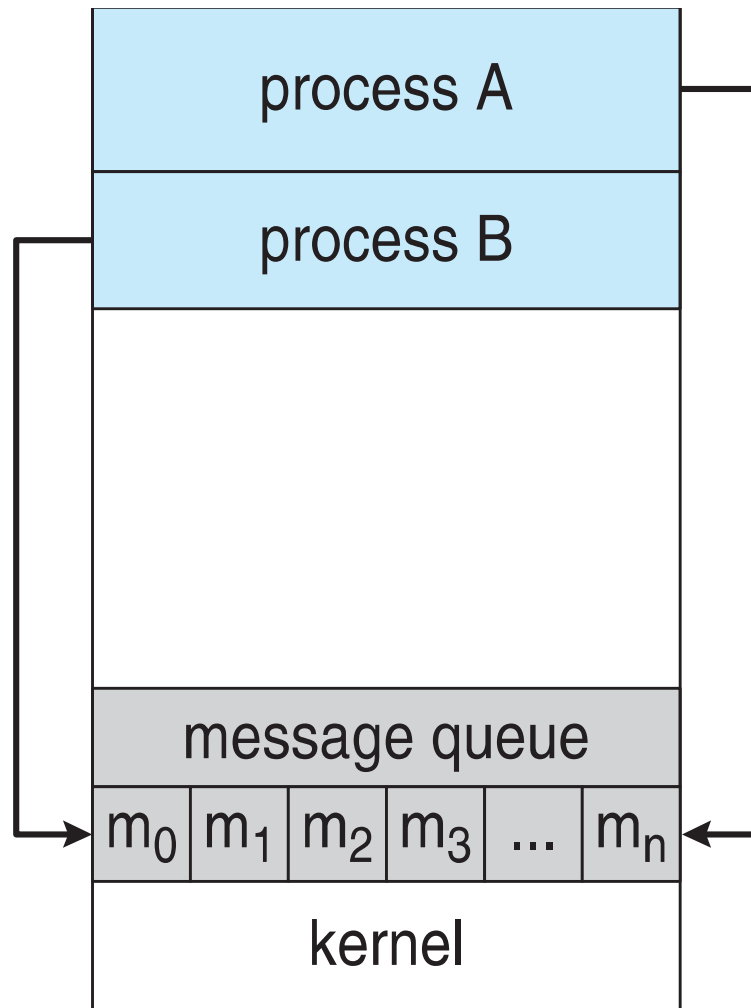
- Processes within a system may be Independent or cooperating
- **Independent process**- that cannot affect or be affected by other processes executing in the system
- **Cooperating process**- It can affect or be affected by other processes executing in the system
- Reasons for process cooperation-
 - Information sharing 
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication**

Models of IPC


- **Shared memory**- A region of memory that is shared by cooperating processes is established. Processes can then exchange information by reading and writing data to the shared region
- **Message passing**- Communication takes place by means of messages exchanged between the cooperating processes.
- **Which ones better?** - Shared memory
- But shared memory suffers from **cache coherency** issues.

Communication Models

(a) Message passing (b) Shared memory



(a)



(b)