

Threads

Moumita Patra
July-Nov 2019
Ref:- Galvin, Gagne

Thread

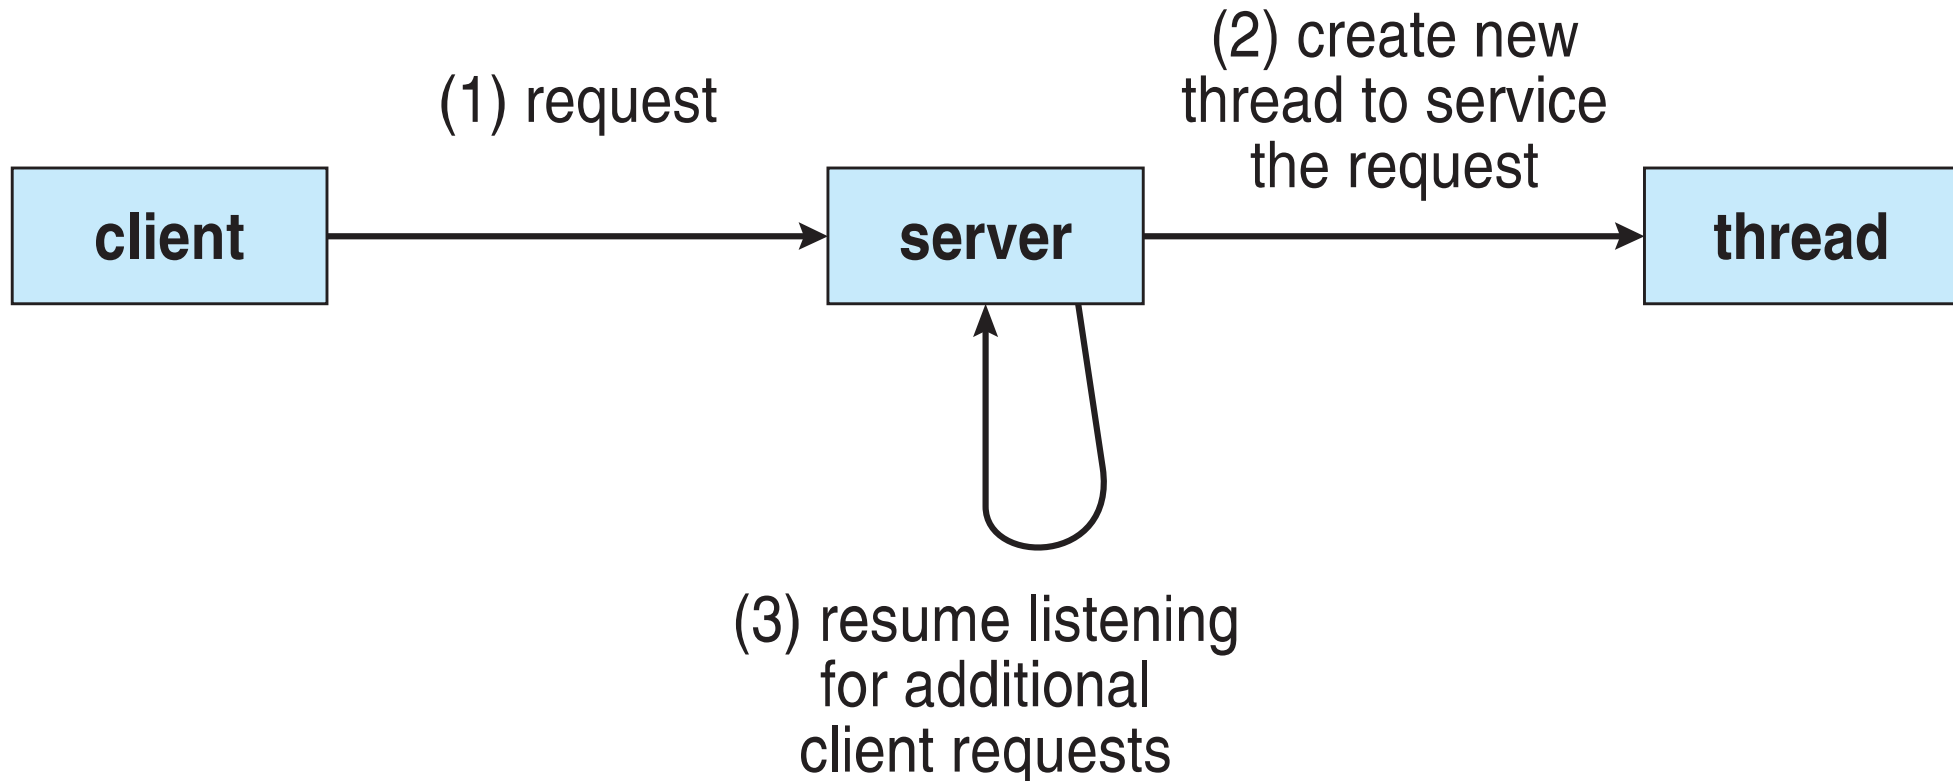
- A thread is a basic unit of CPU utilization.
- It comprises of-
 - Thread ID
 - A program counter
 - A register set
 - A stack
- It shares with other threads belonging to the same process its code section, data section, and other resources, such as open files and signals.

Single and Multithreaded Processes

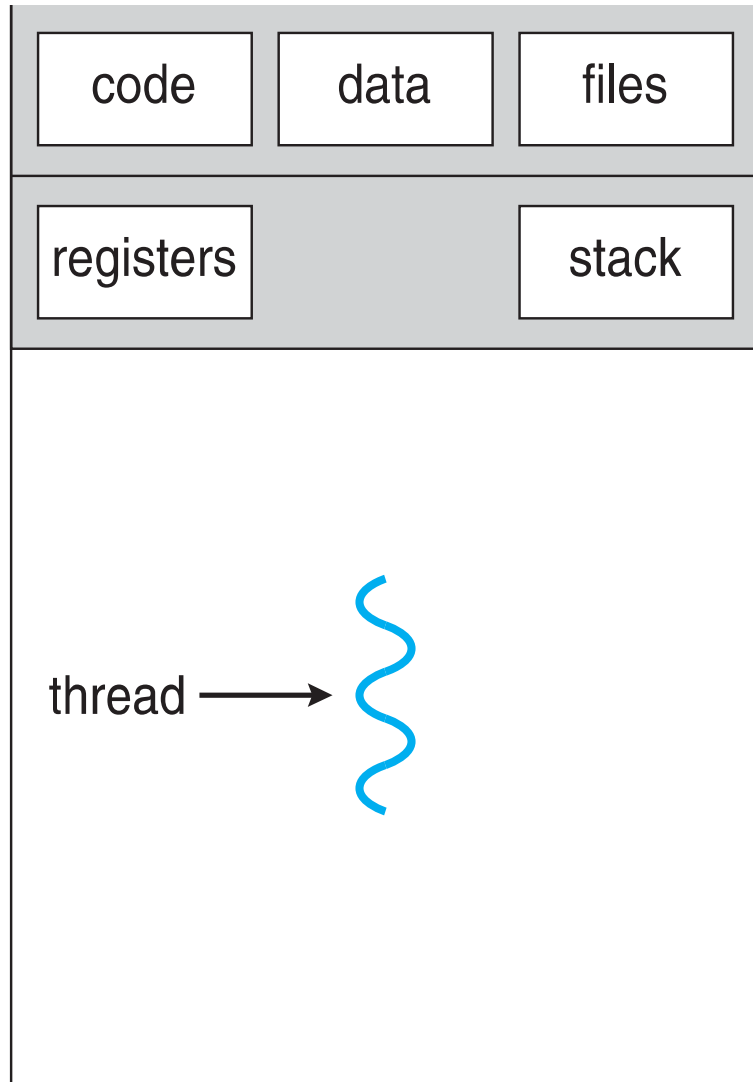
Multithreading

- If a process has multiple threads of control, it can perform more than one task at a time.
- Benefits:-
 - Responsiveness
 - Resource sharing
 - Economy
 - Scalability

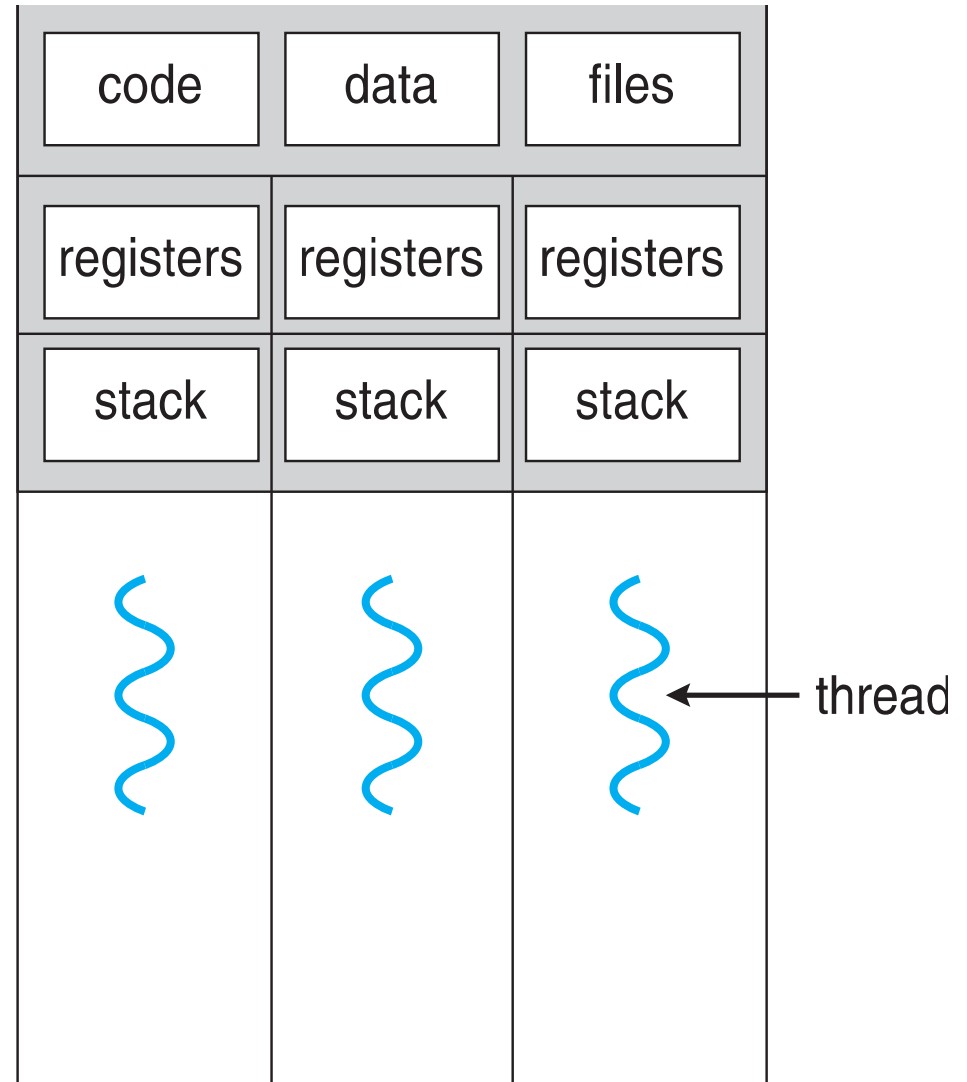
Multithreaded Server Architecture



Single and Multithreaded Processes



single-threaded process



multithreaded process

Multicore Programming

- Multiple computing cores across CPU chips or within CPU chips
- Provides a mechanism for more efficient use of multiple computing cores and improved concurrency.

Multicore Programming


- Challenges in programming for multicore systems:
 - Identifying the tasks
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging

Concurrency and Parallelism


- **Parallelism**- a system that can perform more than one task simultaneously.
- **Concurrency**- supports more than one task by allowing all the tasks to make progress.

Concurrency vs. Parallelism

- Concurrent execution on single-core system:



- Parallelism on a multi-core system:



Multicore Programming (Cont.)

- Types of parallelism
 - **Data parallelism**- distributing subsets of the same data across multiple cores, same operation on each
 - **Task parallelism**- distributing threads across cores, each thread performing unique operation
- As number of threads grows, so does architectural support for threading

Multithreading Models

User Threads and Kernel Threads


- **User threads**- management done by user-level threads library
- Three primary thread libraries:
 - POSIX Pthreads
 - Windows threads
 - Java threads
- Kernel threads- Supported by kernel
- Examples- virtually all general purpose OS, such as:
 - Windows
 - Solaris
 - Linux
 - Mac OS X

Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many

Many-to-One

- Many user-level threads mapped to single kernel thread
- If one thread is blocked, it causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Examples- Solaris Green Threads
- GNU Portable Threads




One-to-One

- Each user-level thread maps to kernel thread
- Creating user-level thread creates a kernel thread
- More concurrency than Many-to-One
- Number of threads per process is sometimes restricted due to overhead


- Examples:-

- Windows
- Linux
- Solaris



Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allow OS to create a sufficient number of kernel threads
- Example- Solaris prior version of 9



Thread Librares

- Thread library provides programmer with API for creating and managing threads
- Two primary ways of implementing-
 - Library entirely in user space
 - Kernel-level library supported by the OS

Pthreads

- A POSIX standard API for thread creation and synchronization
- May be provided as either user-level or kernel-level
- Specification for thread behaviour, not an implementation
- API specifies behaviour of the thread library, implementation is up to development of the library
- Common in UNIX OS

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

```

    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid,&attr,runner,argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}

```

Pthread code for joining 10 threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}
```

Windows Multithreaded C Program (Cont.)

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
}
```


Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS

Java threads may be created by:

```
public interface Runnable
{
    public abstract void run();
}
```

- Extending Thread class
- Implementing the Runnable interface

Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```

Java Multithreaded Program (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>"); }
}
```

Implicit Threading

- For supporting multithreaded applications
- **Implicit threading**- creation and management of threads by compilers and run-time libraries
- Growing in popularity because, as numbers of threads increase, maintaining program correctness becomes more difficult with explicit threads
- Three methods:
 - Thread Pools
 - OpenMP
 - Grand Central Dispatch

Thread Pools

- Create a number of threads in a pool where they wait for work
- When a server receives a request, it awakens a thread from this pool and passes it to the request for service
- Once the thread completes its service, it returns to the pool and awaits for more work.
- If the pool contains no available thread, the server waits until one becomes free.

Benefits-

- Servicing a request with an existing thread is faster than waiting to create a thread
- A thread pool limits the number of threads that exist at any point
Important on systems that cannot support a large number of concurrent threads
- Different strategies can be used for running the task as creation is separate
- Number of threads in the pool can be set heuristically based on factors such as- number of CPUs in the system, amount of physical memory, and the expected number of concurrent client requests

OpenMP

- Set of compiler directives and API for programs written in C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** as blocks of code that may run in parallel.
- Application developers insert compiler directives into their code at parallel regions.
- These directives instruct the openMP run-time library to execute the region in parallel.

Grand Central Dispatch (GCD)

- Technology for Mac OS X and iOS operating systems.
- Combination of extensions to C language, an API, and a run-time library that allows application developers to identify sections of code to run in parallel.
- Extensions are known as **blocks**.
- GCD schedules blocks for run-time execution by placing them on a **dispatch queue**.
- When it removes a block from a queue, it assigns the block to an available thread from the thread pool it manages.

Grand Central Dispatch

Two types of dispatch queue:

- **Serial**- blocks placed on a serial queue are removed in FIFO order
- Each process has its own queue
- Once a block has been removed from the queue, it must complete execution before another block is removed
- **Concurrent**- blocks removed in FIFO order but several blocks may be removed at a time
- There are 3 system-wide concurrent dispatch queues distinguished according to priority- low, default, and high

Threading Issues

- Semantics of `fork()` and `exec()` system calls
- Signal handling
- Thread cancellation
- Thread-local storage
- Scheduler activations

Semantics of fork() and exec()

- Does fork() duplicate only the calling thread or all threads?- some OS have two versions of fork()
- If thread invokes exec(), the program specified in the parameter to exec() will replace the entire process including threads
- Which version of fork() to use depends on the application and when exec() is called

Signals

- **Signals** are used in UNIX systems to notify a process about the occurrence of a particular event
- Two types-
- **Synchronous**- delivered to the same process that performed the operation that caused the signal
- **Asynchronous**- when a signal is generated by an event external to a running process, that process receives the signal asynchronously

Signal Handling

- A **signal handler** is used to process signals
 - Signal is generated by particular event
 - Signal is delivered to a process
 - Signal is handled by either- default or user-defined handler
- Every signal has **default handler** that kernel runs when handling signal
- **User-defined signal handler** can override default

Challenges

Where should a signal be delivered for multi-threaded system?

- Deliver the signal to the thread to which the signal applies
- Deliver the signal to every thread in the process
- Deliver the signal to certain threads in the process
- Assign a specific thread to receive all signals for the process

Thread Cancellation

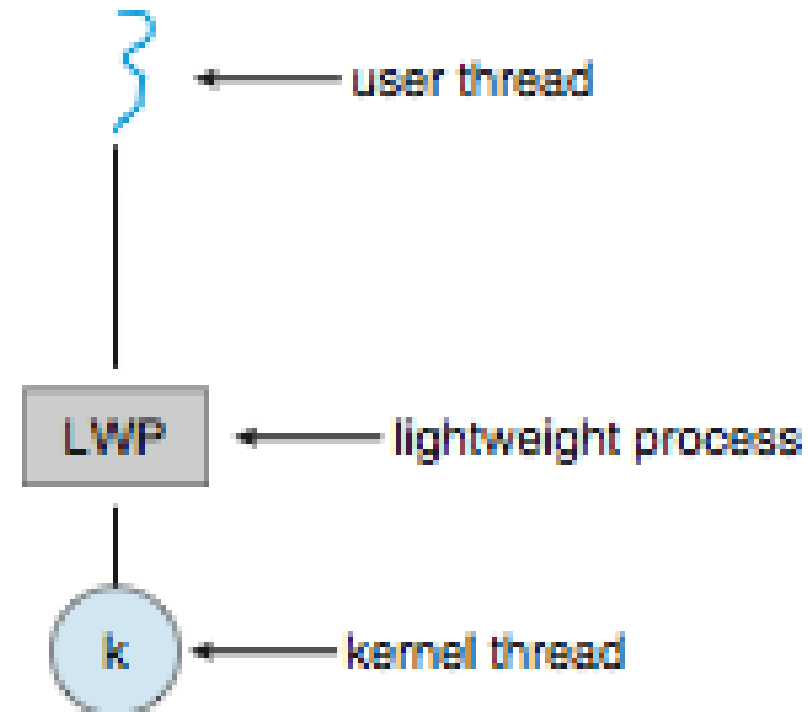
- Terminating a thread before it has completed
- Thread to be canceled is referred as target thread
- Two approaches:
 - **Asynchronous cancellation**- One thread immediately terminates the target thread
 - **Deferred cancellation**- the target thread periodically checks whether it should terminate, allowing it an opportunity to terminate itself in an orderly fashion

Thread-Local Storage (TLS)

- Data that is local to a thread
- TLS allows each thread to have its own copy of data
- Useful when you do not have control over thread creation process, i.e, in case of thread pool
- Different from local variables
 - Local variables are visible only during a single function invocation
 - TLS data are visible across function invocations but are unique to each thread (similar to static data)

Scheduler Activations

- Both many-to-many and two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Typically use an intermediate data structure between user and kernel threads- **lightweight process (LWP)**
 - Appears to be a virtual processor on which process can schedule user thread to run
 - Each LWP attached to kernel thread



Scheduler Activations

- Scheduler activations provide **upcalls**- a communication mechanism from the kernel to the **upcall handler** in the thread library
- This communication allows an application to maintain the correct number of kernel threads

NOTE:

A thread is a flow of control within a process