

# Deadlocks

**Moumita Patra  
Galvin & Gagne  
Jul-Nov 2019**

- A system consists of a finite number of resources to be distributed among a number of competing processes.
- Resources may be partitioned into several types/classes.
- Each resource type has a number of identical instances.
- CPU cycles, files, and I/O devices are examples of resource types.

Each process utilizes a resource as follows:

- Request
- Use
- Release

# Deadlock Characterization

Deadlock can arise if the following conditions hold simultaneously:


- Mutual exclusion
- Hold and wait
- No preemption
- Circular wait




# Resource-Allocation Graph

Deadlocks can be described more precisely in terms of a directed graph called a **system resource-allocation graph**.


- Graph consists of a set of vertices  $V$  and a set of edges  $E$ .
- $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system.
- $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- **request edge** – directed edge  $P_i \rightarrow R_j$
- **assignment edge** – directed edge  $R_j \rightarrow P_i$



# Resource Allocation Graph With Deadlock



# Is there a deadlock here?





- If graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - If only one instance per resource type, then deadlock
  - If several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

Deadlock problem can be dealt in one of three ways:

- Deadlock prevention, deadlock avoidance
- Deadlock detection and recovery
- Ignore the problem altogether and pretend that deadlocks never occur in the system

# Deadlock Prevention

Ensure that at least one of the necessary conditions for deadlock does not hold.

- **Mutual Exclusion**- not required for sharable resources
- **Hold and Wait**-
  - Whenever a process requests a resource, it must not hold any other
  - May release all the resources before requesting additional resources
  - Low resource utilization
  - Starvation is possible

# Deadlock Prevention

- **No Preemption**

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- Preempted resources are added to the list of resources for which the process is waiting.

# Deadlock Prevention

- **Circular Wait**

- Impose a total ordering of all resource types.
- Requirement- each process requests resources in an increasing order of enumeration.

# Deadlock Avoidance

- OS requires additional information in advance concerning which resources a process will request and use during its lifetime.
- Requires that process declare the maximum number of resources of each type that it may need.
- Dynamically examines resource allocation state to ensure that circular-wait never holds.
- Resource allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes.


# Safe State

- A state is **safe** if the system can allocate all resources requested by all processes (upto their stated maximums) without entering a deadlock state.
- System is in **safe state** if there exists a **safe sequence**

*if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the system such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$*

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state





# Avoidance Algorithms


- Single instance of a resource type => use a resource-allocation graph
- Multiple instances of a resource type => use the banker's algorithm

# Resource Allocation Graph Scheme


- **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_i$  may request resource  $R_j$ ; represented by a dashed line
- Claim edge converts to request edge when a process requests a resource
- Request edge is converted to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge

Resources must be claimed *apriori* in the system

# Resource-Allocation Graph



# Unsafe State In Resource-Allocation Graph



# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# Banker's Algorithm

- Multiple instances of resources
- Each process must declare maximum number of each resource type required apriori.
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

# Data Structures for the Banker's Algorithm

$n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If  $\text{available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $\text{Max}[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task



$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

# Safety Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively. Initialize:

**Work** = **Available**

**Finish** [ $i$ ] = **false** for  $i = 0, 1, \dots, n-1$

2. Find an  $i$  such that both:

(a) **Finish** [ $i$ ] == **false**

(b) **Need** <sub>$i$</sub>  ≤ **Work**

If no such  $i$  exists, go to step 4

3. **Work** = **Work** + **Allocation** <sub>$i$</sub>

**Finish** [ $i$ ] = **true**

go to step 2

4. If **Finish** [ $i$ ] == **true** for all  $i$ , then the system is in a safe state



# Resource-Request Algorithm for Process $P_i$

**$Request_i$**  = request vector for process  $P_i$ . If  **$Request_i[j] = k$**  then process  $P_i$  wants  **$k$**  instances of resource type  $R_j$

1. If  **$Request_i \leq Need_i$** , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  **$Request_i \leq Available$** , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

If safe  $\Rightarrow$  the resources are allocated to  $P_i$

If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored


# Deadlock Detection

- Allow system to enter deadlock state
- Detect
- Recover

# Single Instance of Each Resource Type


- Maintain **wait-for** graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph

# Resource-Allocation Graph and Wait-for Graph



(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph

# Data Structures

- **Available:** A vector of length  $m$  indicates the number of available resources of each type
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If **Request**  $[i][j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

1. Let ***Work*** and ***Finish*** be vectors of length ***m*** and ***n***, respectively. Initialize:

(a) ***Work = Available***

(b) For ***i = 1, 2, ..., n***, if ***Allocation<sub>i</sub> ≠ 0***, then ***Finish[i] = false***; otherwise, ***Finish[i] = true***

2. Find an index ***i*** such that both:

(a) ***Finish[i] == false***

(b) ***Request<sub>i</sub> ≤ Work***

If no such ***i*** exists, go to step 4

3.  **$Work = Work + Allocation_i$**   
 **$Finish[i] = true$**   
go to step 2

4. If  **$Finish[i] == false$** , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  **$Finish[i] == false$** , then  $P_i$  is deadlocked

**Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state**

# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will be affected by the deadlock?
    - => one for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.



# Recovery from Deadlock

# Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?

# Resource Preemption

Successively preempt some processes and give the resources to other processes until the deadlock cycle is broken.

Issues:

- **Selecting a victim**- select a process to terminate. Which one?
- **Rollback**- return to some safe state. What to do with rolled back process?

restart process from that state

- **Starvation**- same process may always be picked as victim.  
Include number of rollback in cost factor (solution).

A process can be picked as a victim only a finite number of times.