

Congestion Control in TCP

Dr. Manas Khatua
Assistant Professor
Dept. of CSE, IIT Guwahati
E-mail: manaskhatua@iitg.ac.in

Principles of Congestion Control





- We have discussed: **reliable data transfer** service in the face of packet loss
 - such loss typically results from the **overflowing of router buffers** as the network becomes congested
- **Packet retransmission** treats a symptom of network congestion **but not** the cause of network congestion

congestion:


- “too many sources sending too much data too fast for **network** to handle”
- different from flow control!
- manifestations:
 - **lost packets** (buffer overflow at routers)
 - **long delays** (queueing in router buffers)
- a top-10 problem!

Causes/Cost of Congestion : scenario 1

- ❖ two senders, two receivers
- ❖ one router, **infinite buffers**
- ❖ output **link capacity: R**
- ❖ no error recovery (i.e. no retransmission)



- ❖ max. per-connection throughput: $R/2$
- ❖ no matter how high their sending rates!
- ❖ as the router buffer is shared among them.




- ❖ large delays as arrival rate λ_{in} approaches capacity

Conclusion:

- while operating at an aggregate throughput of near R
 - may be **ideal from a throughput standpoint**,
 - but, it is **far from ideal from a delay standpoint!**
- Even in this (extremely) idealized scenario
 - “cost” of a congested network
 - large **queuing delays** are experienced as the **packet arrival rate** nears the **link capacity**.
 - **assuming** that: (i) the connections operate at these sending rates for an infinite period of time, (ii) there is an infinite amount of buffering available
→ the above delay between source and destination **becomes infinite**

Causes/costs of Congestion: scenario 2

- ❖ one router, *finite* buffers
- ❖ sender *retransmission* of timed-out packet
 - application-layer input = application-layer output: $\lambda_{in} = \lambda_{out}$
 - transport-layer input includes *retransmissions* : $\lambda'_{in} \geq \lambda_{in}$
- ❖ Sending rate: λ_{in} ; Offered load: λ'_{in}





Cont...

Assume idealization: perfect knowledge with sender

- ❖ sender sends only when router buffers available

❖ so, no loss $\rightarrow \lambda_{in} = \lambda_{out}$




Cont...

Assume Idealization:

know when loss occur

- ❖ sender only re-sends if packet *known* to be lost





Cont...

Assume Idealization:

know when loss occur



- ❖ sender only re-sends if packet *known* to be lost



Cont...

Realistic: *duplicates*


- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender **times out prematurely**, sending **two** copies, both of which are delivered



Cont...

Realistic: duplicates

- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender **times out prematurely**, sending **two** copies, both of which are delivered



“costs” of congestion:


- ❖ more work (retrans) for given “goodput”
- ❖ unneeded retransmissions: link carries multiple copies of pkt
 - decreasing goodput

Causes/costs of congestion: scenario 3

- ❖ four senders
- ❖ multihop paths
- ❖ timeout/retransmit
- ❖ **Overlapping** paths
- ❖ all have same value of λ_{in}


Q: what happens if all λ_{in} and λ_{in}' increase ?

A: as red λ_{in}' increases, all arriving blue pkts at upper queue (R1) are dropped; blue throughput $\rightarrow 0$



Cont...

- ❖ For extremely small values of λ_{in} , buffer overflows are rare
 - the throughput approximately equals the offered load
- ❖ For slightly larger values of λ_{in} , overflows are still rare
 - the corresponding throughput is also larger,
- ❖ Thus, for small values of λ_{in} , an increase in λ_{in} results in an increase in λ_{out}



As red λ_{in}' increases, all arriving blue pkts at upper queue (in R1) are dropped, as R1 will give priority to red pkts;
So, blue throughput $\rightarrow 0$

another “cost” of congestion:

- ❖ when packet is dropped, any “upstream transmission capacity” used for that packet was wasted! (e.g. work by R4 in above figure)

Congestion v/s Flow Control



- TCP **cannot ignore** the congestion in network (at the intermediate points) as it wants to provide end-to-end reliability
- The use of **flow control** in TCP **cannot avoid congestion** in **intermediate routers** because
 - a router may receive data from more than one sender
 - Flow control is for individual TCP sender
 - There is no congestion at the either end
 - there may be congestion in the middle.

Approaches to Congestion Control



two broad approaches towards congestion control:

end-to-end congestion control

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system who observed loss, delay
- ❖ approach taken by TCP
- ❖ suitable in datagram approach

network-assisted congestion control

- ❖ routers/switches provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM ABR)
 - explicit rate for sender to send at
 - Direct feedback: sent from a network router to the sender
 - Indirect feedback: router marks a field in a packet flowing from sender to receiver
- ❖ suitable for virtual-circuit approach

TCP Congestion control



- Basic approach:
 - each sender limit the rate at which it sends traffic into its connection
 - set the rate as a function of perceived network congestion.
- perceives less congestion along the path → increases its send rate
- perceives huge congestion along the path → reduces its send rate
- It should not aggressively send segments to the network
- It can not be very conservative, either, sending a small number of segments in each time interval

Cont...

- **Questions** need to answer:
 - How does a TCP sender **limit the rate** at which it sends traffic into its connection?
 - How does a TCP sender **perceive that there is congestion** on the path between itself and the destination?
 - What **congestion control algorithm** should the sender use to change its send rate as a function of perceived end-to-end congestion?

Answer of 1st Question:

- To control the number of segments to transmit, TCP uses another variable called Congestion Window (*cwnd*)
- Actually, the *cwnd* variable and the *rwnd* variable (used for flow control) together define the **size of the send window** in TCP
 - Actual send window size = $\min(rwnd, cwnd)$
- The constraint above limits the amount of unacknowledged data at the sender and therefore indirectly limits the sender's send rate.


Answer of 2nd Question:

- TCP sender uses the occurrence of two events as **signs of congestion**:
 - time-out
 - 3 duplicate ACKs

Answer of 3rd Question:

- There **exist many congestion control algorithm** for adjusting the value of **cwnd** based upon end-to-end congestion
 - Default/basic approach
- **Modified TCP with congestion control** algorithms
 - **Tahoe TCP**: both signs of occurrence are treated equally
 - **Reno TCP**: both signs of occurrence are treated differently
 - **New Reno TCP**: TCP checks to see if more than one segment is lost in the current window when 3 duplicate ACKs arrive

TCP Congestion Control: details



TCP sending rate:

- ❖ *roughly*: send cwnd bytes, wait RTT for ACKs, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- ❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteACKed} \leq \text{cwnd}$$


- ❖ **cwnd** is dynamic, function of perceived network congestion

TCP congestion control algo has three components:


- ❖ *slow start*
- ❖ *congestion avoidance*
- ❖ *fast recovery*

Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
 - initially **cwnd** = 1 MSS (maximum-sized segments)
 - double **cwnd** every RTT
 - done by incrementing **cwnd** for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast
- ❖ This process results in a doubling of the sending rate every RTT.



When growth ends?

- **1st case**, a **loss event** indicated by a **timeout** 
 - Indicates congestion
 - **cwnd** sets to 1 MSS
 - begins the **slow start** process anew.
 - **ssthresh** (**slow start threshold**) sets to $cwnd/2$.
- **2nd case**, when the value of **cwnd** equals **ssthresh**,
 - TCP transitions into **congestion avoidance** state
 - **cwnd** grows linearly
- **3rd case**, if **3 duplicate ACKs** are detected,
 - dupACKs indicate network capable of delivering some segments
 - TCP performs a fast retransmit and enters **fast recovery** state
 - **ssthresh** sets to $cwnd/2$.
 - **cwnd** sets to $ssthresh + 3 \text{ MSS}$.
 - **cwnd** grows linearly


Slow-start strategy is **slower** in the case of **delayed ACK**.

If two segments are ACKed **cumulatively**, the size of the **cwnd** **increases by 1, not 2**. With one ACK for every two segments, the growth is a power of 1.5, but **still exponential**

Congestion Avoidance

- On entry to this state, the value of *cwnd* is approx **half its value** when congestion was last encountered
- To **avoid congestion** before it happens, we must **slow down the exponential growth** of *cwnd*
- the **additive** phase begins.
- If **3 dupACKs** are detected at this state,
 - TCP performs a fast retransmit and enters the **fast recovery state**
- If **timeout** occurs at this state
 - TCP enters into **slow start**

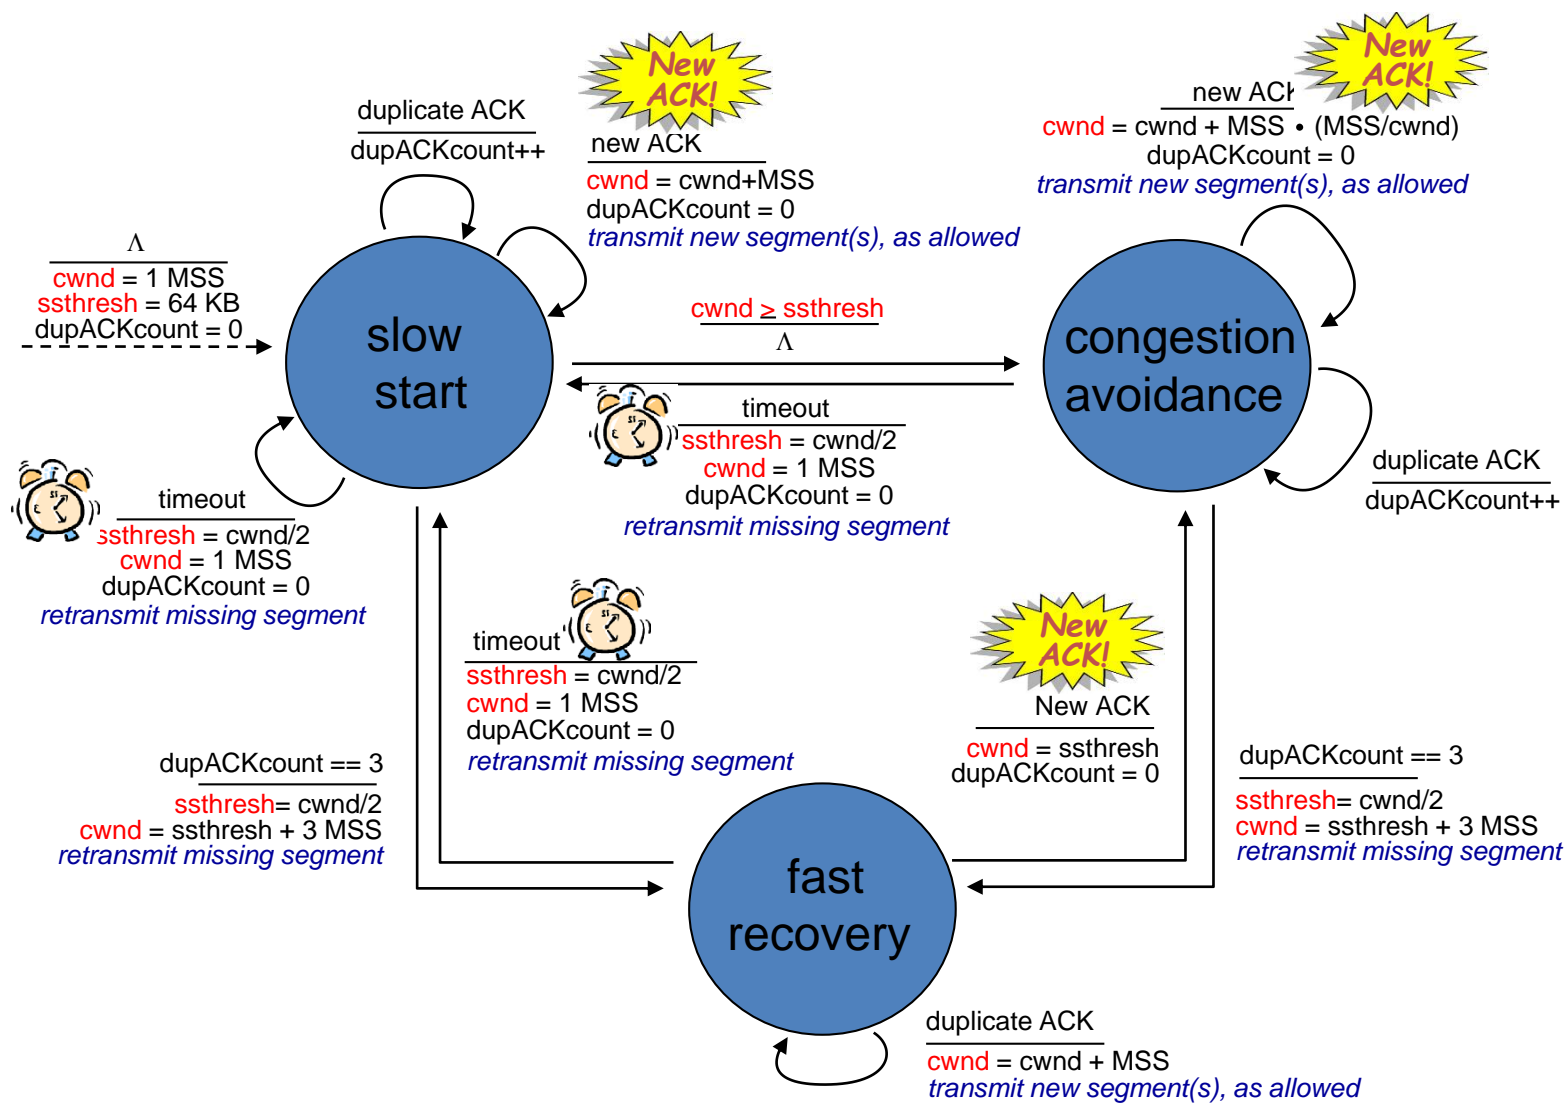
If a new ACK arrives, $cwnd = cwnd + MSS \cdot (MSS/cwnd)$



Fast Recovery

- this algorithm is also an **additive increase**, but it **starts when 3 duplicate ACK arrives**
- If a **duplicate ACK** arrives (after the 3 duplicate ACK which triggers the recovery)
 - $cwnd = cwnd + (1/ cwnd)$
- If **timeout** occurs, TCP moves back to **slow start** state
- If any **new ACK** arrives, TCP moves back to **congestion avoidance** state
- This state is recommended, but not mandatory in TCP

FSM of TCP Congestion Control




Different Versions

- **TCP Tahoe**
 - signs of congestion occurrence (time-out, 3 duplicate ACK) are treated equally
 - uses only *slow start* and *congestion avoidance* states
- **TCP Reno**
 - signs of congestion occurrence (time-out, 3 duplicate ACK) are treated differently
 - three states in FSM: *slow start*, *congestion avoidance*, *fast recovery*
- **TCP New Reno**
 - It differs from RENO in that it *doesn't exit fast-recovery until* all the data which was outstanding at the time it entered fast recovery is ACKed.
 - *It is most common today*
- **TCP Vegas**
 - variations of the Reno algorithm
 - attempts to avoid congestion while maintaining good throughput
 - The basic idea of Vegas is to
 - (1) *detect congestion* in the routers between source and destination *before packet loss occurs*, and
 - (2) *lower the rate linearly* when this imminent *packet loss is detected*.

Additive Increase Multiplicative Decrease

- TCP congestion control is often referred to as AIMD form of congestion control.
- ❖ *approach*: sender increases **transmission rate** (**window** size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase window by 1 MSS every RTT until loss detected
 - *multiplicative decrease*: cut window in half after loss


AIMD saw-toothed behavior: probing for bandwidth



TCP Throughput

- What the **average throughput** of a **long-lived TCP connection** would be?
- we'll **ignore the slow-start** phases **that occur after timeout** events as these phases are typically very short.
- the **rate at which TCP sends data** is a function of **cwnd** and **current RTT**
 - $\text{Rate} = \text{cwnd} / \text{RTT}$
- Let, $\text{cwnd} = W$ **when a loss event occurs**.

If we ignore
slow-start then



- Assume that RTT and W are **approximately constant over the duration** of the connection (i.e. in steady-state)
 - the TCP transmission rate ranges from $(W / 2 \text{ RTT})$ to (W / RTT)
- So, the **average throughput** of a connection = $\frac{1}{2} ((W / 2 \text{ RTT}) + (W / \text{RTT})) = 0.75 * (W / \text{RTT})$

TCP over “High-Bandwidth” path

- Example of high speed TCP needed in present era:
 - 1500 byte segments, 100ms RTT,
 - We want 10 Gbps throughput
- So, using previous formula --> it requires $W = 83,333$ in-flight segments
- What would happen the case of loss?
- throughput in terms of segment loss probability, L [Mathis 1997]:


$$\text{TCP throughput} = \frac{1.22 \cdot \text{MSS}}{\text{RTT} \sqrt{L}}$$

→ to achieve 10 Gbps throughput, need a loss rate of $L = 2 \cdot 10^{-10}$
it means very small loss rate!

- new versions of TCP for high-speed

TCP Fairness

fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K



Thanks!

Content of this PPT are taken from:

- 1) **Computer Networks: A Top Down Approach**, by J.F. Kurose and K.W. Ross, 6th Eds, 2013, Pearson Education.
- 2) **Data Communications and Networking**, by B. A. Forouzan , 5th Eds, 2012, McGraw-Hill.
- 3) **Chapter 3 : Transport Layer**, PowerPoint slides of “Computer Networking: A Top Down Approach”, 6th Eds, J.F. Kurose, K.W. Ross