


Message Passing

Message Passing

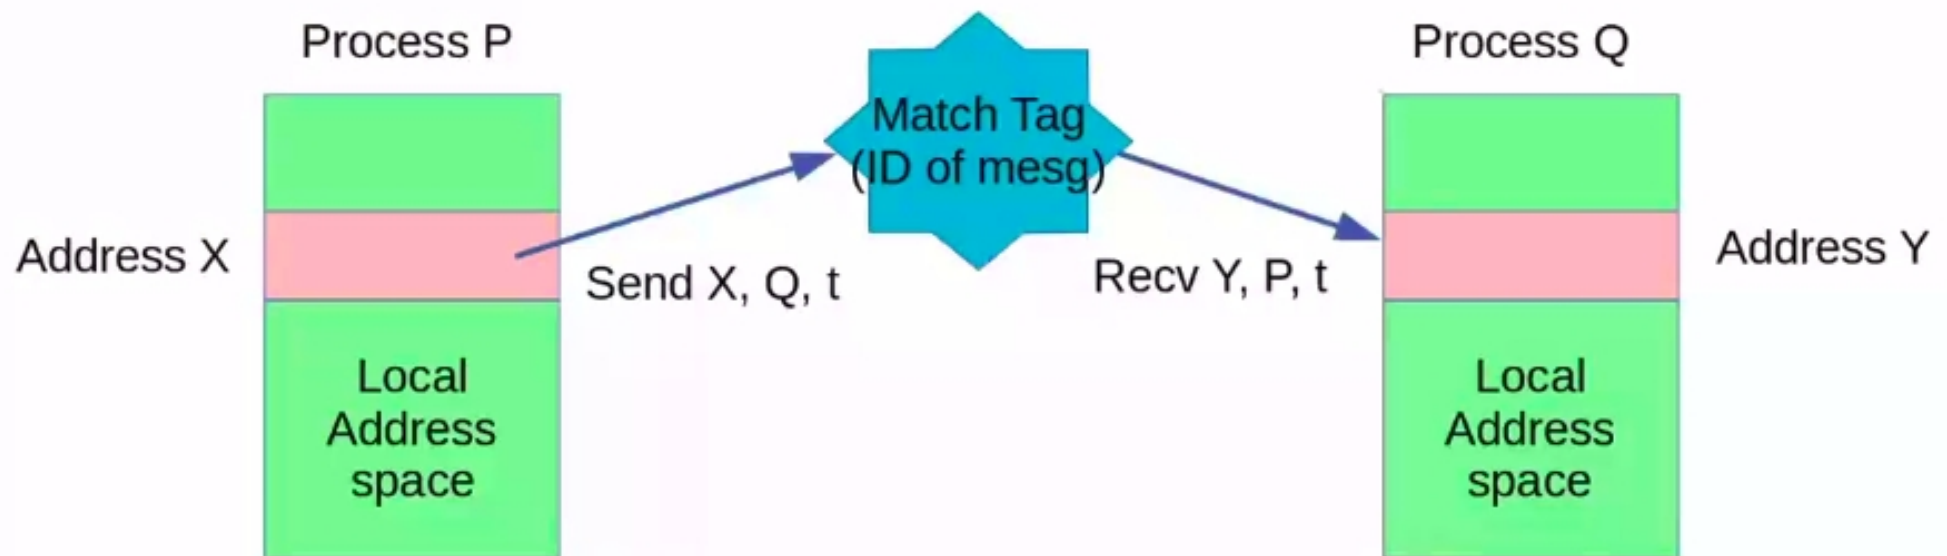
- Message passing architectures have computers as building blocks (processor + I/O system) + provides communication between processors as explicit I/O operations
- Takes help of OS and library calls for actual implementation
- Unlike local area networks, this has more tight integration
- This is processor to processor communication
- Common user-level operations are variants of SEND/RECEIVE
- SEND: specifies local buffer that has to be transmitted + to whom (receiving process)
- RECEIVE: specifies sending process + local buffer into which the received data has to be kept

Message Passing Architectures




- **Cannot directly access memory on another node**
- **IBM SP-2, Intel Paragon, clusters of PCs ("beowulf")**

Message Passing




Message Passing

- Send and Receive accomplish a pairwise synchronisation event and perform a memory-to-memory copy, as each process provides its local address
- Possible variants:
 - Whether Send
 - completes when receiver completes
 - When is send buffer available for reuse
 - When new requests can be accepted
 - Whether receiver
 - Waits till matching send occurs Or simply post the request 
- Message passing was used in earlier programming languages like: CSP, Occam and OS functions like sockets

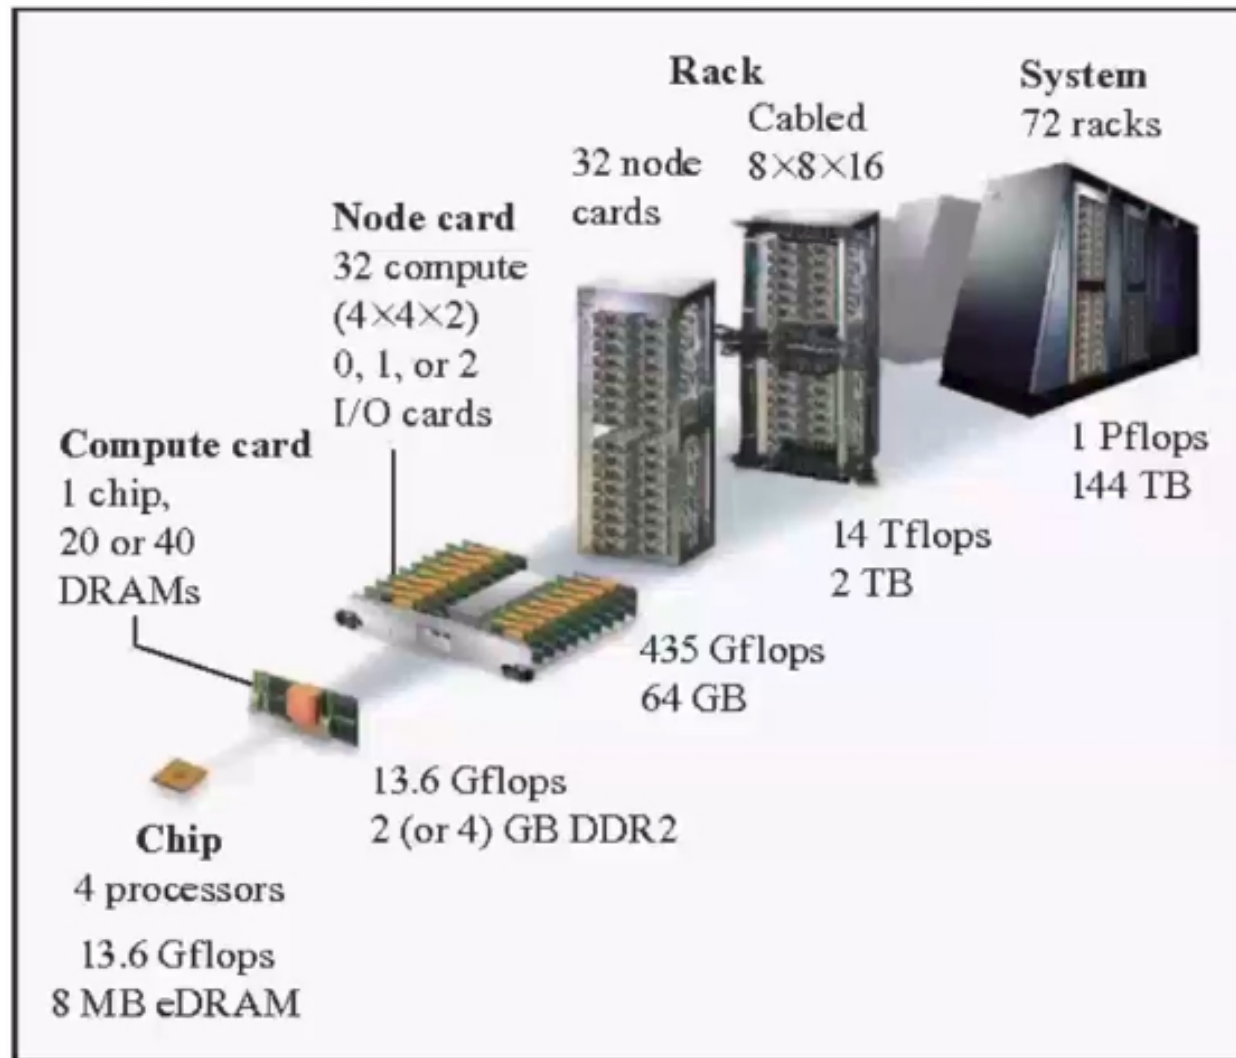
Message passing




- Earlier message passing machines used point-to-point network
- Used FIFO to send data. Network topology was important
 - Sender wrote to the link and Receiver read from the link
 - Sender will block until receiver reads the value = called synchronous message passing
- Some used DMA + dedicated processor for send/receive
 - Allowed non-blocking send
- Later developments allowed message to go to any destination via other nodes. Network level routing protocols
 - Store-and-forward networks
 - Latency depended on number of hops traversed



IBM Blue Gene/P supercomputer



Infiniband network



Blue Gene Router logic

